

# Distributed Systems

ECE428

Lecture 9

*Adopted from Spring 2021*

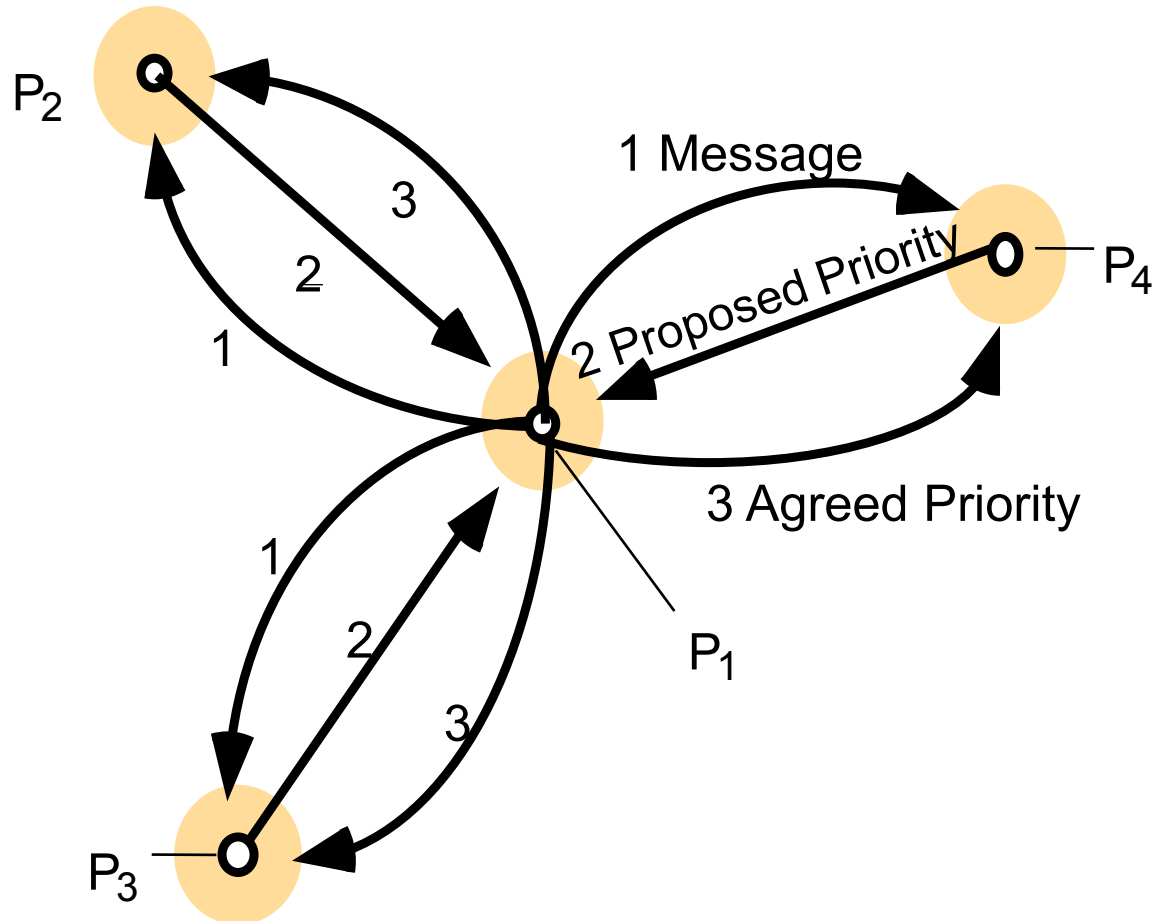
# Today's agenda

- Wrap up Multicast
  - Chapter 15.4
  - Tree-based multicast and Gossip
- Mutual Exclusion
  - Chapter 15.2

# Recap: Ordered Multicast

- FIFO ordering: If correct process issues  $\text{multicast}(g, m)$  and then  $\text{multicast}(g, m')$ , then every correct process that delivers  $m'$  will have already delivered  $m$ .
- Causal ordering: If  $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$  then any correct process that delivers  $m'$  will have already delivered  $m$ .
  - Note that  $\rightarrow$  counts multicast messages delivered to the application, rather than all network messages.
- Total ordering: If a correct process delivers message  $m$  before  $m'$ , then any other correct process that delivers  $m'$  will have already delivered  $m$ .

# ISIS algorithm for total ordering



Proposed Priority: *higher than all priorities proposed by the process and agreed priorities received by the process so far.*

Agreed Priority: *Maximum of all proposed priority for the message*

# Proof of total order with ISIS

- Consider messages,  $m_1$  and  $m_2$ , and two processes,  $p$  and  $p'$ .
- Suppose that  $p$  delivers  $m_1$  before  $m_2$ .
- When  $p$  delivers  $m_1$ , it is at head of the queue.  $m_2$  is either:
  - Already in  $p$ 's queue, and deliverable, so
    - $\text{Final\_priority}(m_1) < \text{Final\_priority}(m_2)$
  - Already in  $p$ 's queue, and not deliverable, so
    - $\text{Final\_priority}(m_1) < \text{Proposed\_priority}(m_2) \leq \text{Final\_priority}(m_2)$
  - Not yet in  $p$ 's queue:
    - same as above, since proposed priority  $>$  priority of any delivered message
- Suppose  $p'$  delivers  $m_2$  before  $m_1$ , by the same argument:
  - $\text{Final\_priority}(m_2) < \text{Final\_priority}(m_1)$
  - Contradiction!

# Ordered Multicast

- FIFO ordering
  - If a correct process issues  $\text{multicast}(g, m)$  and then  $\text{multicast}(g, m')$ , then every correct process that delivers  $m'$  will have already delivered  $m$ .
- Causal ordering
  - If  $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$  then any correct process that delivers  $m'$  will have already delivered  $m$ .
  - Note that  $\rightarrow$  counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
  - If a correct process delivers message  $m$  before  $m'$  then any other correct process that delivers  $m'$  will have already delivered  $m$ .

# Implementing causal order multicast

- Similar to FIFO Multicast
  - What you send with a message differs.
  - Updating rules differ.
- Each receiver maintains a vector of per-sender sequence numbers (integers)
  - Processes  $P_1$  through  $P_N$ .
  - $P_i$  maintains a vector of sequence numbers  $P_i[1 \dots N]$  (initially all zeroes).
  - $P_i[j]$  is the latest sequence number  $P_i$  has received from  $P_j$ .

*... these are NOT vector logical clocks!*


# Implementing causal order multicast

- *CO-multicast*( $g, m$ ) at  $P_j$ :
  - set  $P_j[j] = P_j[j] + 1$
  - piggyback entire vector  $P_j[1 \dots N]$  with  $m$ .
  - B-multicast*( $g, \{m, P_j[1 \dots N]\}$ )
- On *B-deliver*( $\{m, V[1 \dots N]\}$ ) at  $P_i$  from  $P_j$ : If  $P_i$  receives a multicast from  $P_j$  with sequence vector  $V[1 \dots N]$ , buffer it until both conditions are true:
  1. This message is next one  $P_i$  is expecting from  $P_j$ , i.e.,  
 $V[j] = P_i[j] + 1$
  2. All multicasts, anywhere in the group, which happened before  $m$  have been received at  $P_i$ , i.e.,  
For all  $k \neq j$ :  $V[k] \leq P_i[k]$

When these conditions satisfied, CO-deliver( $m$ ), and set  $P_i[j] = V[j]$




# Causal order multicast execution




Self-deliveries omitted for simplicity.

# Causal order multicast execution



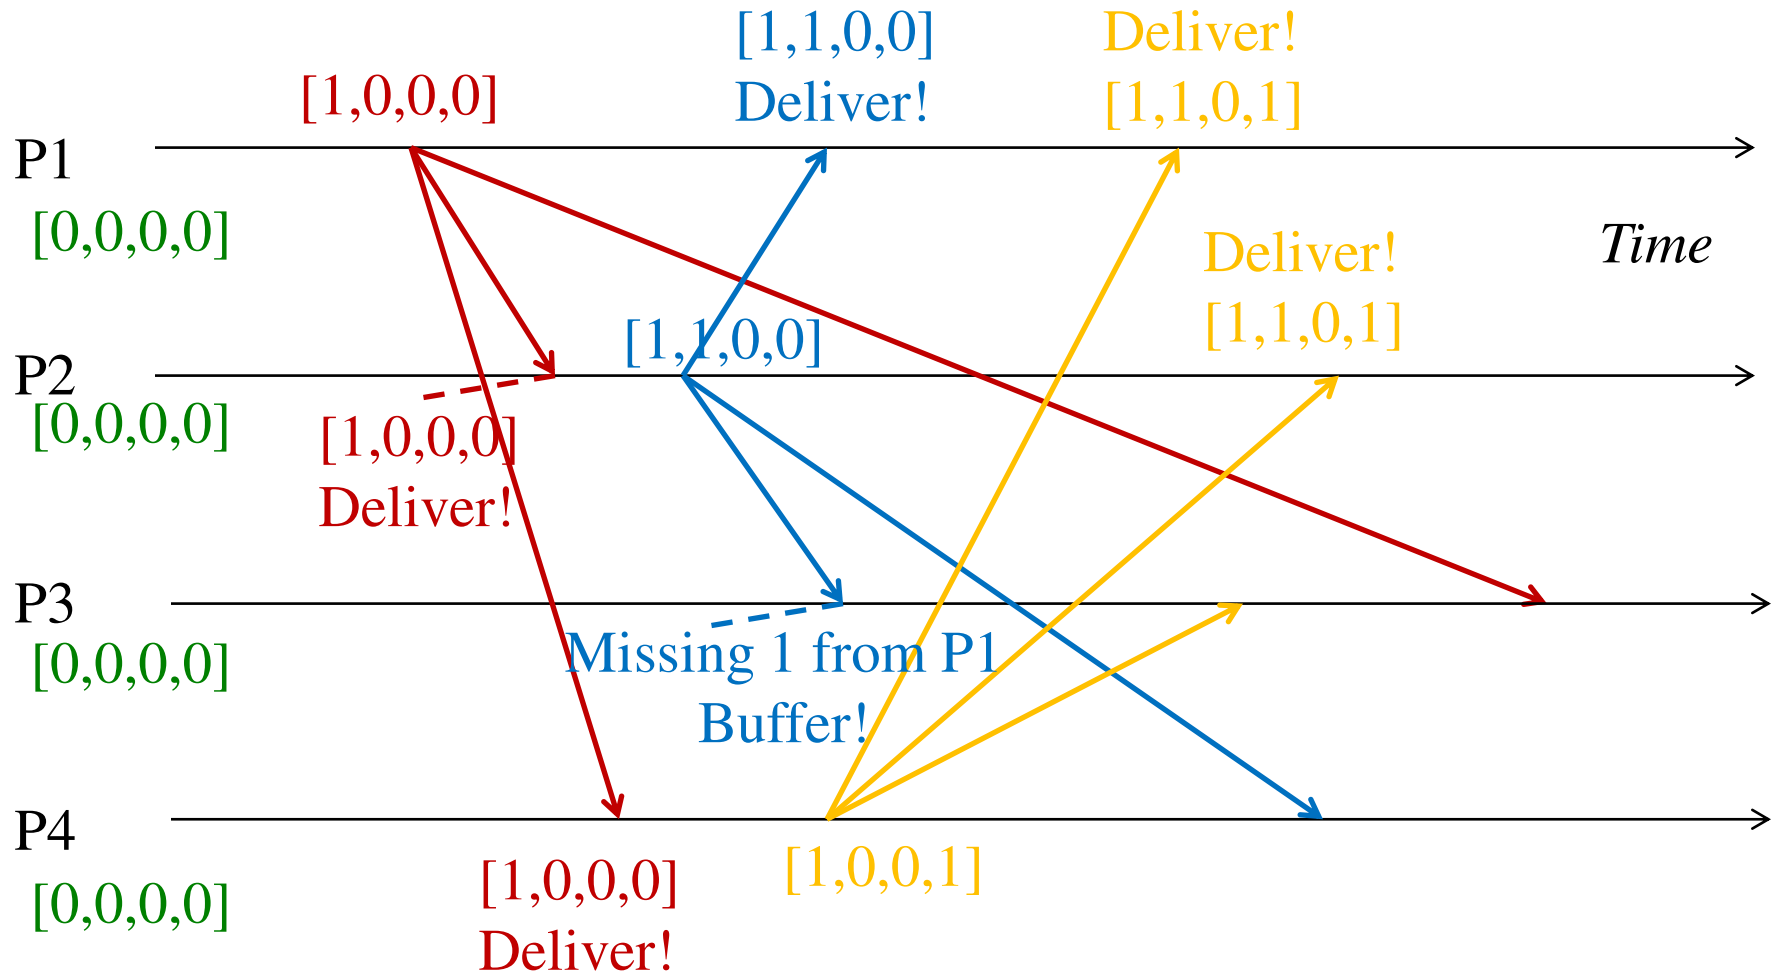
Self-deliveries omitted for simplicity.

# Causal order multicast execution



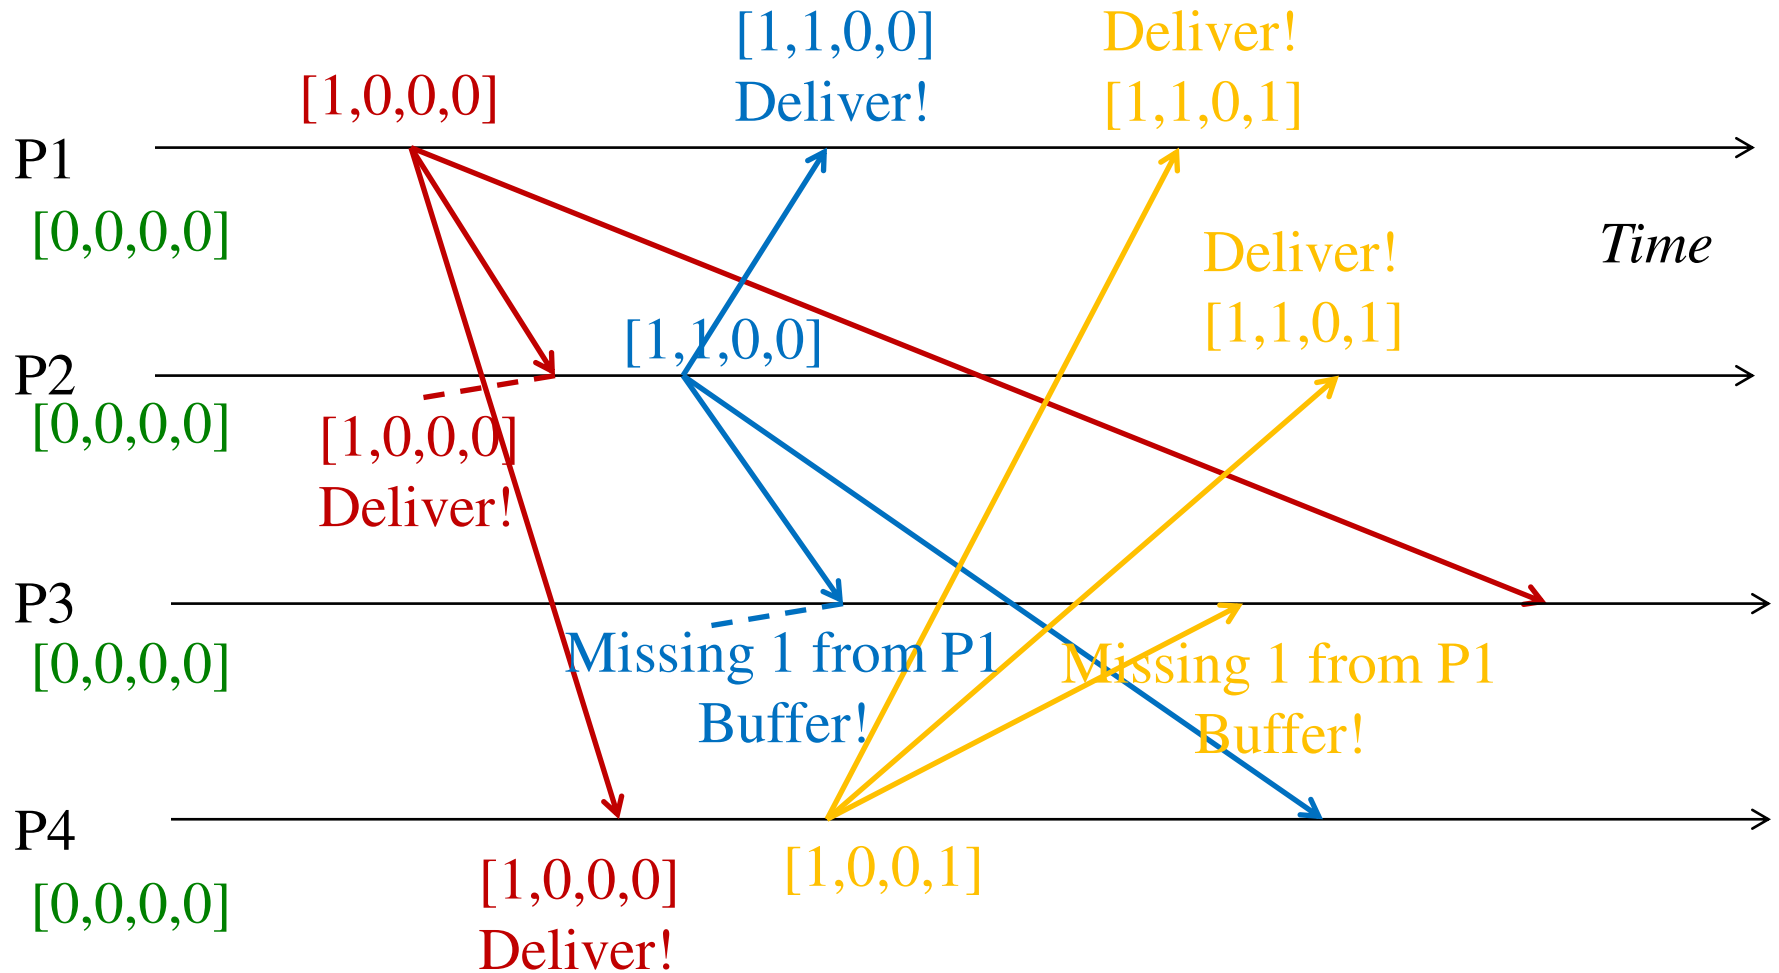
Self-deliveries omitted for simplicity.

# Causal order multicast execution



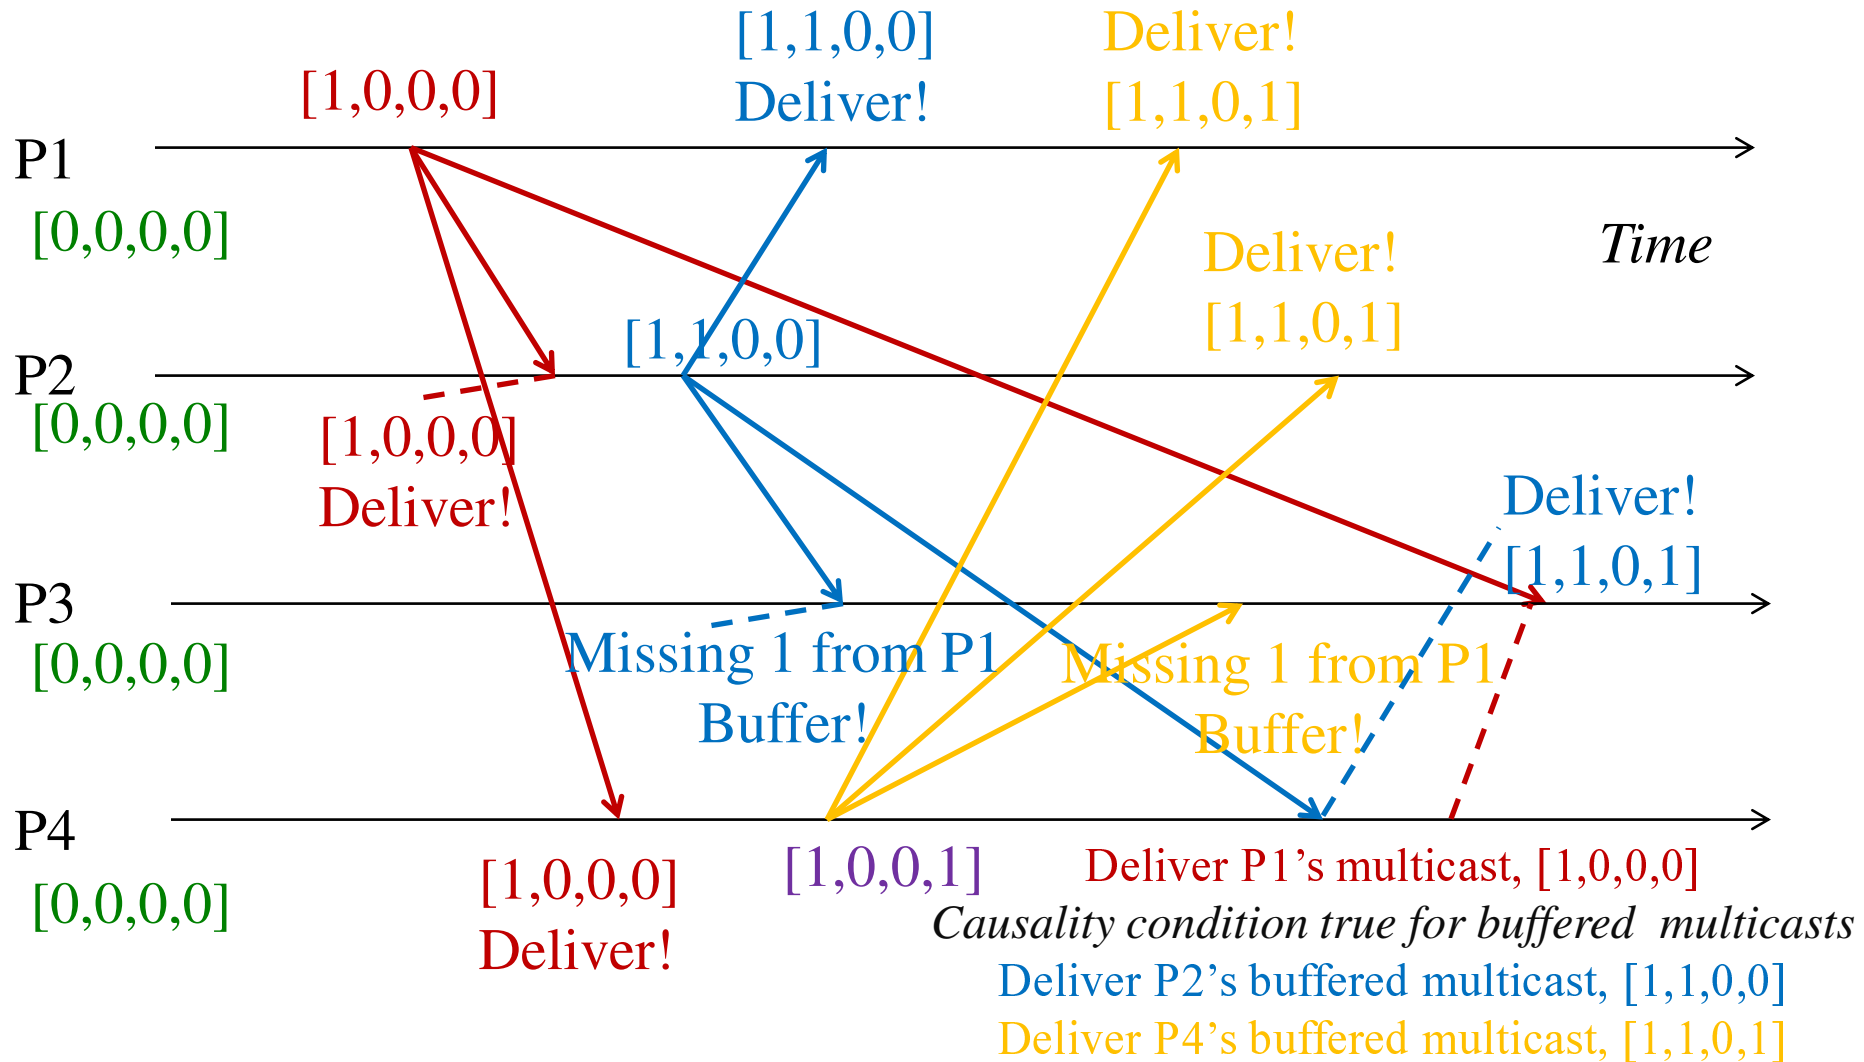
Self-deliveries omitted for simplicity.

# Causal order multicast execution



Self-deliveries omitted for simplicity.

# Causal order multicast execution



# Causal order multicast implementation

- Only looks at multicast messages delivered to the application.
- Ignores causality created due to other network messages.

# Ordered Multicast

- FIFO ordering
  - If a correct process issues  $\text{multicast}(g, m)$  and then  $\text{multicast}(g, m')$ , then every correct process that delivers  $m'$  will have already delivered  $m$ .
- Causal ordering
  - If  $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$  then any correct process that delivers  $m'$  will have already delivered  $m$ .
  - Note that  $\rightarrow$  counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
  - If a correct process delivers message  $m$  before  $m'$ , then any other correct process that delivers  $m'$  will have already delivered  $m$ .








# More efficient multicast mechanisms


- Our focus so far has been on the application-level semantics of multicast.
- *What are some of the more efficient underlying mechanisms for a B-multicast?*

# B-Multicast

Sender




# B-Multicast using unicast sends




# B-Multicast using unicast sends

Sender

Closer look at physical network paths.




# B-Multicast using unicast sends



# B-Multicast using unicast sends

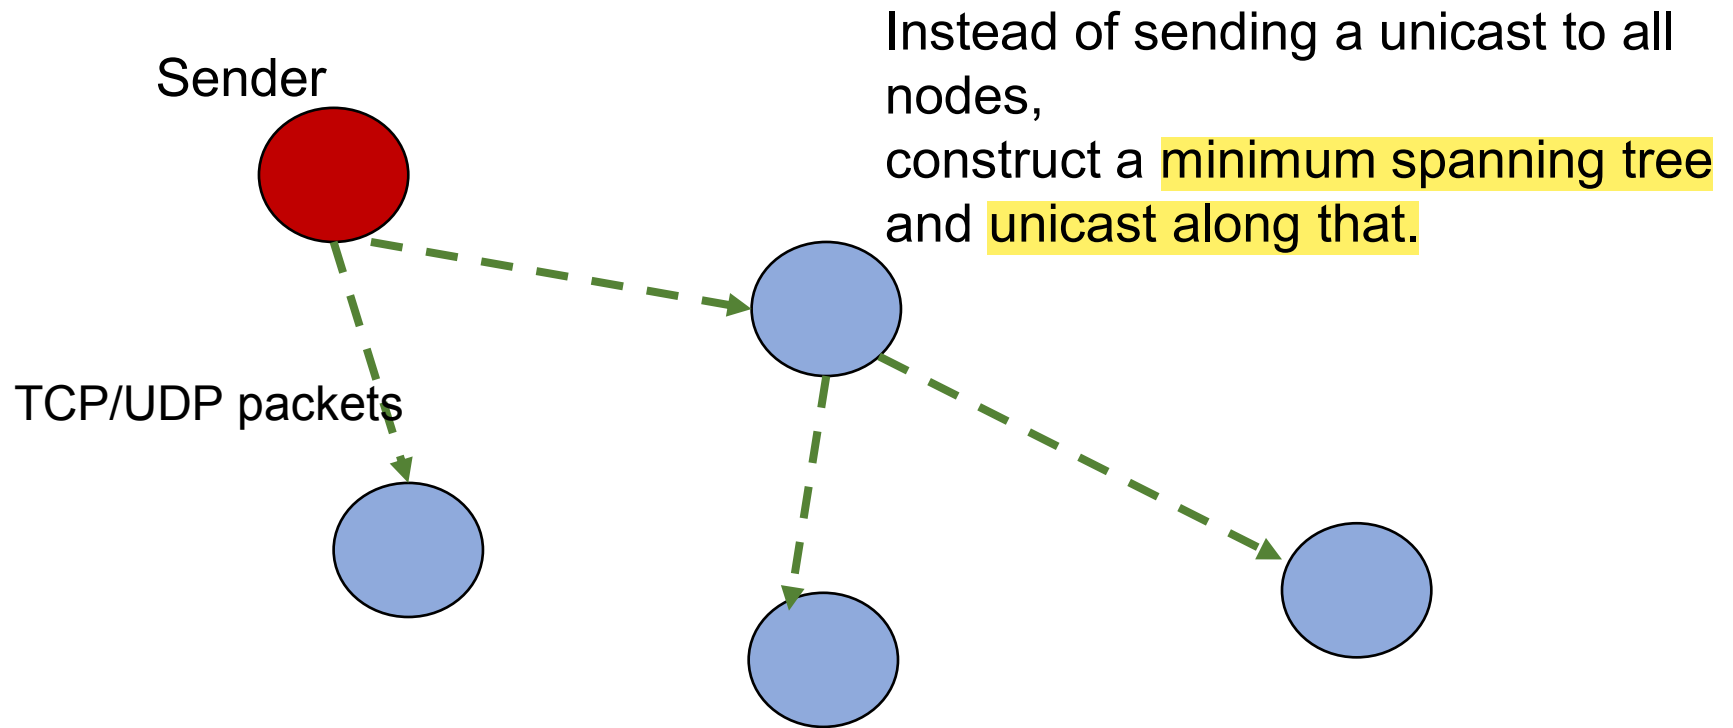
Similar redundancy when individual nodes also act as routers (e.g. wireless sensor networks).

Sender




How do we reduce the overhead?

# Tree-based multicast



# Tree-based multicast

A process does not directly send messages to *all* other processes in the group.






# Tree-based multicast

A process does not directly send messages to *all* other processes in the group.

Sender


It sends a message to only a subset of processes.

*Closer look at the physical network.*




# Tree-based multicast

Sender




Also possible to construct a tree that includes network routers. IP multicast!

# Tree-based multicast




# Third approach: Gossip

Transmit to  $b$  random targets.




# Third approach: Gossip

Transmit to  $b$  random targets.



# Third approach: Gossip

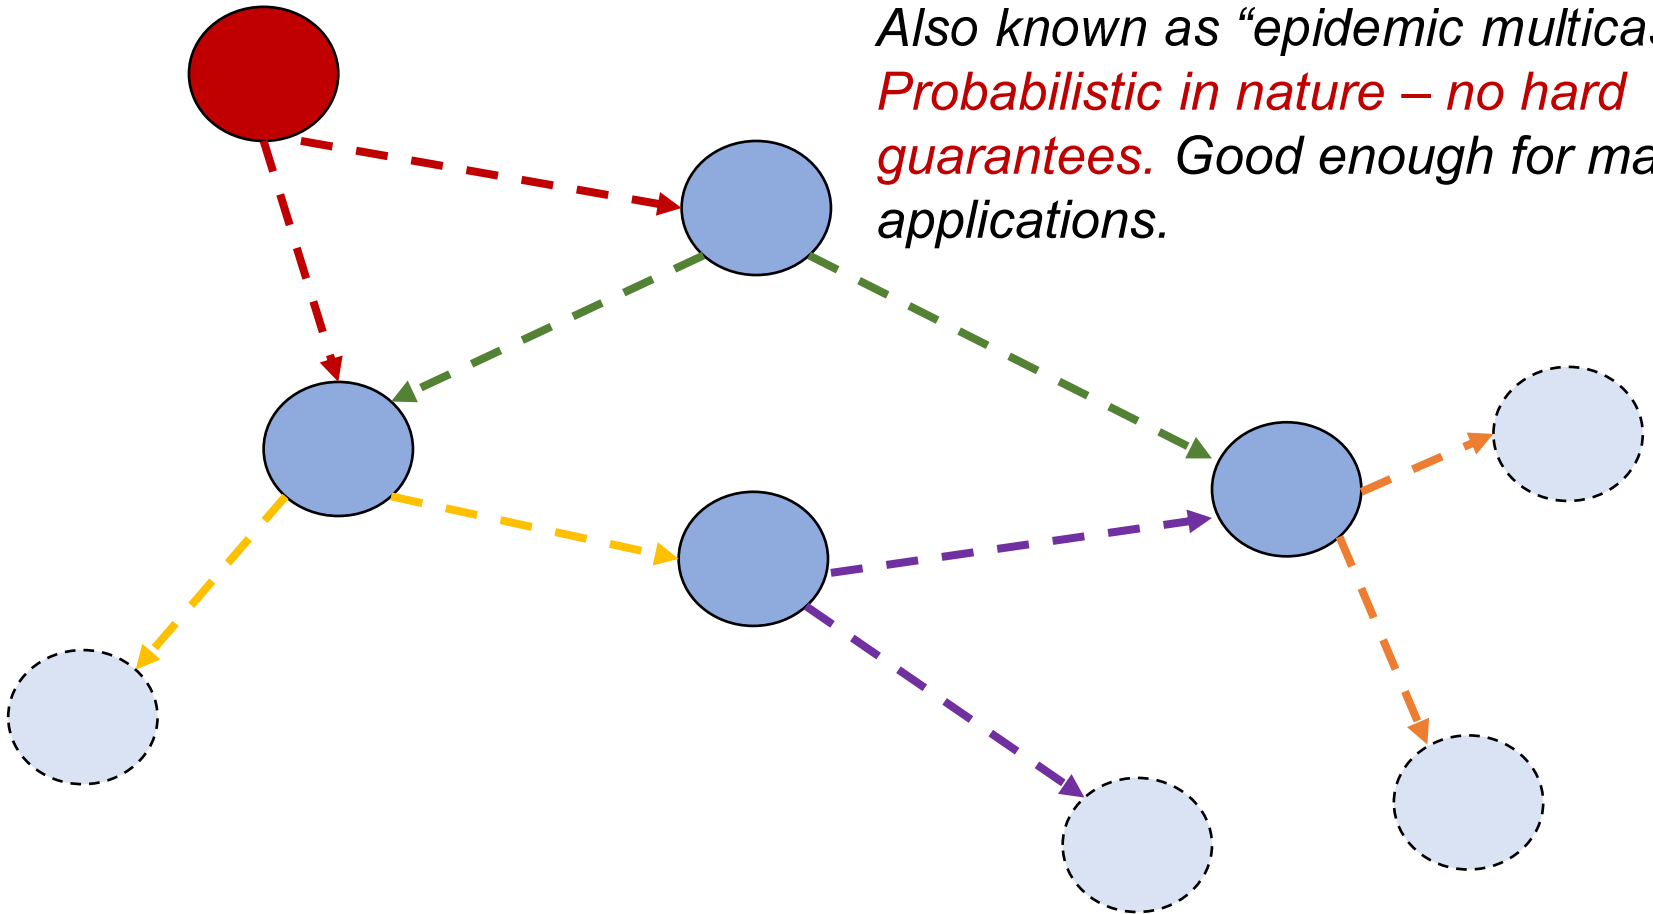
Transmit to  $b$  random targets.



# Third approach: Gossip

No “tree-construction” overhead.  
More efficient than unicasting to all  
receivers.


*Also known as “epidemic multicast”.  
**Probabilistic in nature – no hard  
guarantees.** Good enough for many  
applications.*



# Third approach: Gossip

Used in many real-world systems:

- Facebook's distributed datastore uses it to determine group membership and failures.
- Bitcoin uses it to exchange transaction information between nodes.





# Multicast Summary

- Multicast is important for applications in distributed systems.
- Applications may have different requirements:
  - Basic
  - Reliable
  - Ordering: FIFO, Causal, Total
  - Combinations of the above.
- Underlying mechanisms to spread the information:
  - Unicast to all receivers.
  - Tree-based multicast, and gossip: sender unicasts messages to only a subset of other processes, and they spread the message further.
  - Gossip is more scalable and more robust to failures.

# Today's agenda

- Wrap up Multicast
  - Chapter 15.4
  - Tree-based multicast and Gossip
- Mutual Exclusion
  - Chapter 15.2
- Goal: reason about ways in which different processes in a distributed system can safely manipulate shared resources.

# Why Mutual Exclusion?

- **Bank's Servers in the Cloud:** Two of your customers make simultaneous deposits of \$1,000 into your bank account, each from a separate ATM.
  - Both ATMs read initial amount of \$10,000 in your account concurrently from the bank's cloud server
  - Both ATMs add \$1,000 to this amount (locally at ATM)
  - Both write the final amount to the server
  - **What's wrong?**

# Why Mutual Exclusion?

- **Bank's Servers in the Cloud:** Two of your customers make simultaneous deposits of \$1,000 into your bank account, each from a separate ATM.
  - Both ATMs read initial amount of \$10,000 in your account concurrently from the bank's cloud server
  - Both ATMs add \$1,000 to this amount (locally at ATM)
  - Both write the final amount to the server
  - **You lost \$1,000!**
- **The ATMs need *mutually exclusive* access to your account entry at the server**
  - or, mutually exclusive access to executing the code that modifies the account entry.

# More uses of mutual exclusion

- Distributed file systems
  - Locking of files and directories
- Accessing objects in a safe and consistent way
  - Ensure at most one server has access to object at any point of time
- In industry
  - Chubby is Google's locking service

# Problem statement for mutual excls.

- *Critical Section Problem*:
  - Piece of code (at all processes) for which we need to ensure there is at most one process executing it at any point of time.
- Each process calls three functions
  - **enter()** in order to enter the critical section (CS)
  - **AccessResource()** to run the critical section code
  - **exit()** to exit the critical section

# Our bank example

ATM1:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

ATM2:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

# Mutual exclusion for a single OS

- If all processes are running in one OS on the same machine (or VM):
  - Semaphores
  - Mutexes
  - Condition variables
  - Monitors
  - ...



# Processes sharing an OS:

## Semaphores

- Semaphore == an integer that can only be accessed via two special functions
- Semaphore S=1; // Max number of allowed accessors.

**wait(S)** (or P(S) or down(S)):

```
while(1) { // each execution of the while loop is atomic
    if (S > 0) {
```

```
        S--;
        break;
```

```
    }
```

```
}
```

**signal(S)** (or V(S) or up(s)):

```
S++; // atomic
```

**exit()**

Atomic operations are supported via hardware instructions such as compare-and-swap, test-and-set, etc.

# Our bank example

ATM1:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

ATM2:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

# Our bank example

Semaphore S=1; // shared

ATM1:

```
wait(S); //enter
    // AccessResource()
obtain bank amount;
add in deposit;
update bank amount;
    // AccessResource() end
signal(S); // exit
```

ATM2:

```
wait(S); //enter
    // AccessResource()
obtain bank amount;
add in deposit;
update bank amount;
    // AccessResource() end
signal(S); // exit
```

# Mutual exclusion in distributed systems

- Processes communicating by passing messages.
- Cannot share variables like semaphores!
- *How do we support mutual exclusion in a distributed system?*

# Mutual exclusion in distributed systems

- Our focus today: Classical algorithms for mutual exclusion in distributed systems.
  - Central server algorithm
  - Ring-based algorithm
  - Ricart-Agrawala Algorithm
  - Maekawa Algorithm

# Mutual Exclusion Requirements

- Need to guarantee 3 properties:
  - **Safety** (essential):
    - At most one process executes in CS (Critical Section) at any time.
  - **Liveness** (essential):
    - Every request for a CS is granted eventually.
  - **Ordering** (desirable):
    - Requests are granted in the order they were made.

# System Model

- Each pair of processes is connected by reliable channels (such as TCP).
- Messages sent on a channel are eventually delivered to recipient, and in FIFO order.
- Processes do not fail.
  - Fault-tolerant variants exist in literature.

# Mutual exclusion in distributed systems

- Our focus today: Classical algorithms for mutual exclusion in distributed systems.
  - Central server algorithm
  - Ring-based algorithm
  - Ricart-Agrawala Algorithm
  - Maekawa Algorithm



# Central Server Algorithm

- Elect a central server (or leader)
- Leader keeps
  - A queue of waiting requests from processes who wish to access the CS
  - A special token which allows its holder to access CS
- Actions of any process in group:
  - **enter()**
    - Send a request to leader
    - Wait for token from leader
  - **exit()**
    - Send back token to leader

# Central Server Algorithm

- Leader Actions:
  - On receiving a request from process  $P_i$ 
    - if (leader has token)
      - Send token to  $P_i$
    - else
      - Add  $P_i$  to queue
  - On receiving a token from process  $P_i$ 
    - if (queue is not empty)
      - Dequeue head of queue (say  $P_j$ ), send that process the token
    - else
      - Retain token

# Analysis of Central Algorithm

- Safety – at most one process in CS
  - Exactly one token
- Liveness – every request for CS granted eventually
  - With  $N$  processes in system, queue has at most  $N$  processes
  - If each process exits CS eventually and no failures, liveness guaranteed
- Ordering:
  - FIFO ordering guaranteed in order of requests received at leader
  - Not in the order in which requests were sent or the order in which processes enter CS!

# Analysis of Central Algorithm

- Safety – at most one process in CS
  - Exactly one token
- Liveness – every request for CS granted eventually
  - With  $N$  processes in system, queue has at most  $N$  processes
  - If each process exits CS eventually and no failures, liveness guaranteed
- Ordering:
  - FIFO ordering guaranteed in order of requests received at leader
  - Not in the order in which requests were sent or the order in which processes enter CS!

# To be continued in next class

- Metrics for analyzing performance of mutual exclusion algorithms.
- Other algorithms for mutual exclusion in distributed systems.
  - Central server algorithm
  - Ring-based algorithm
  - Ricart-Agrawala Algorithm
  - Maekawa Algorithm