

Distributed Systems

ECE428

Lecture 12

Adopted from Spring 2021

While we wait...

- Think:
 - Which algorithms we have studied so far **require a leader?**
 - How can such a leader be elected in a distributed system?
 - What are the safety and liveness conditions for leader election?

Today's agenda

- Leader Election
 - Chapter 15.3
- Goal:
 - What is leader election in distributed systems?
 - How do we elect a leader?
 - To what extent can we handle failures when electing a leader?

Why Election?

- Example: Your Bank account details are replicated at a few servers, but one of these servers is responsible for receiving all reads and writes, i.e., it is the leader among the replicas
 - What if servers disagree about who the leader is?
 - What if there are two leaders per customer?
 - What if the leader crashes?

Each of the above scenarios leads to inconsistency

More motivating examples

- The root server in a group of NTP servers.
- The master in Berkeley algorithm for clock synchronization.
- In the sequencer-based algorithm for total ordering of multicasts, the “sequencer” = leader.
- The central server in the “central server algorithm” for mutual exclusion.
- Other systems that need leader election: Apache Zookeeper, Google’s Chubby.

Leader Election Problem

- Among processes, elect *Leader* to undertake special tasks
 - And *let everyone know* in the group about this Leader
- What happens when a leader fails (crashes)
 - Some process detects this (using a Failure Detector!)
 - Then what?
- Focus of this lecture: **Election algorithm**. Its goal:
 1. **Elect one leader only among the non-faulty processes**
 2. All non-faulty processes agree on who is the leader

Calling for an Election

- Any process can call for an election.
- A process can call for at most one election at a time.
- Multiple processes are allowed to call an election simultaneously.
 - All of them together must yield only a single leader
- The result of an election should not depend on which process calls for it.

Election Problem, Formally

- A run of the election algorithm must always guarantee:
 - **Safety:** For all non-faulty processes p :
 - p has elected:
 - q : a particular non-faulty process with the *best or unique attribute value* or Null (None)
 - **Liveness:** For all election runs:
 - election run terminates
 - & for all non-faulty processes p : p 's elected is not Null
- At the end of the election protocol, the non-faulty process with the *best (highest) election attribute* value is elected.
 - Common attribute: the leader has the highest id
 - Other attribute examples: leader has the highest IP address, or fastest CPU, or most disk space, or most files, most battery etc.

System Model (Assumptions)


- N processes.
- Messages are eventually delivered.
- Failures may occur during the election protocol.
- Each process has a unique id.
 - Each process has a unique attribute (based on which Leader is elected).
 - If two processes have the same attribute, combine the attribute with the process id to break ties.

Classical Election Algorithms

- Ring election algorithm
- Bully algorithm

Ring Election Algorithm


- N processes are organized in a logical ring
 - All messages are sent clockwise around the ring.



Ring Election Protocol (basic version)


- When P_i start election
 - send election message with P_i 's $\langle attr_i, i \rangle$ to ring successor.
- When P_j receives message (election, $\langle attr_x, x \rangle$) from predecessor
 - If $(attr_x, x) > (attr_j, j)$:
 - forward message (election, $\langle attr_x, x \rangle$) to successor
 - If $(attr_x, x) < (attr_j, j)$
 - send (election, $\langle attr_j, j \rangle$) to successor
 - If $(attr_x, x) = (attr_j, j)$: P_j is the elected leader (why?)
 - send elected message containing P_j 's id
- elected message forwarded along the ring until it reaches the leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election: Example



Goal: Elect highest id process as leader

Ring Election: Example




Goal: Elect highest id process as leader

Ring Election Protocol (basic version)

- When P_i start election
 - send election message with P_i 's $\langle attr_i, i \rangle$ to ring successor.
- When P_j receives message (election, $\langle attr_x, x \rangle$) from predecessor
 - If $(attr_x, x) > (attr_j, j)$:
 - forward message (election, $\langle attr_x, x \rangle$) to successor
 - If $(attr_x, x) < (attr_j, j)$
 - send (election, $\langle attr_j, j \rangle$) to successor
 - If $(attr_x, x) = (attr_j, j)$: P_j is the elected leader (why?)
 - send elected message containing P_j 's id
- elected message forwarded along the ring until it reaches the leader

What happens when multiple processes call for an election?

Ring Election: Example




Election: 80 sent twice.

Elected: 80 also sent twice.

Ring Election Protocol [Chang & Roberts'79]

- When P_i start election
 - send election message with P_i 's $\langle \text{attr}_i, i \rangle$ to ring successor.
 - set state to participating
- When P_j receives message (election, $\langle \text{attr}_x, x \rangle$) from predecessor
 - If $(\text{attr}_x, x) > (\text{attr}_j, j)$:
 - forward message (election, $\langle \text{attr}_x, x \rangle$) to successor
 - set state to participating
 - If $(\text{attr}_x, x) < (\text{attr}_j, j)$
 - If (not participating):
 - send (election, $\langle \text{attr}_j, j \rangle$) to successor
 - set state to participating
 - If $(\text{attr}_x, x) = (\text{attr}_j, j)$: P_j is the elected leader (why?)
 - send elected message containing P_j 's id
- elected message forwarded along the ring until it reaches the leader
 - Set state to not participating when an elected message is received.

Ring Election: Example




Election: 80 and Elected: 80
sent only once.

Performance Analysis


- Let's assume no failures occur during the election protocol itself, and there are N processes.
- Let's also assume that only one process initiates the algorithm
- Bandwidth usage: Total number of messages sent.
- Turnaround time: The number of serialized message transmission times between the initiation and termination of a single run of the algorithm.

Worst-case




When the initiator is
the ring successor of
the would-be leader.

Worst-case




- (N-1) messages for **Election** message to get from N6 to N80.
- N messages for **Election** message to circulate around ring without message being changed.
- N messages for **Elected** message to circulate around the ring
- No. of messages: $(3N-1)$
- **Turnaround time: $(3N-1)$ message transmission times**

Best-case



When the initiator is
the would-be leader.

Best-case



When the initiator is the would-be leader.

No. of messages: $2N$


Turnaround time:
 $2N$ message transmission times

Performance Analysis

- Let's assume no failures occur during the election protocol itself, and there are N processes.
- Let's also assume that only one process initiates the algorithm
- Bandwidth usage (total number of messages)
 - $O(N)$: Worst case = $3N - 1$; Best case = $2N$.
- $O(N)$ turnaround time.

Performance Analysis


- Let's assume no failures occur during the election protocol itself, and there are N processes.
- When each process initiates the algorithm?
 - $O(N)$ messages in best-case.



- N election messages generated at the start of algorithm.
- Only one survives, and completes a full round.
 - $N-1$ messages.
- One round for the elected message
 - N messages.
- Total: $3N - 1$ messages

Performance Analysis

- Let's assume no failures occur during the election protocol itself, and there are N processes.
- When each process initiates the algorithm?
 - $O(N)$ messages in best-case.
 - $O(N^2)$ in worst-case.



- N election messages generated at the start of algorithm.
- $N - 1$ survive the next time step.
- $N - 2$ survive the next time step.
-


Performance Analysis

- Let's assume no failures occur during the election protocol itself, and there are N processes.
- When each process initiates the algorithm?
 - $O(N)$ messages in best-case.
 - $O(N^2)$ messages in worst-case.
 - $O(N)$ turnaround time.

Correctness

- Assuming no process fails.
- Safety:
 - Process with highest attribute elected by all nodes.
- Liveness:
 - Election completes within $3N - 1$ message transmission times.

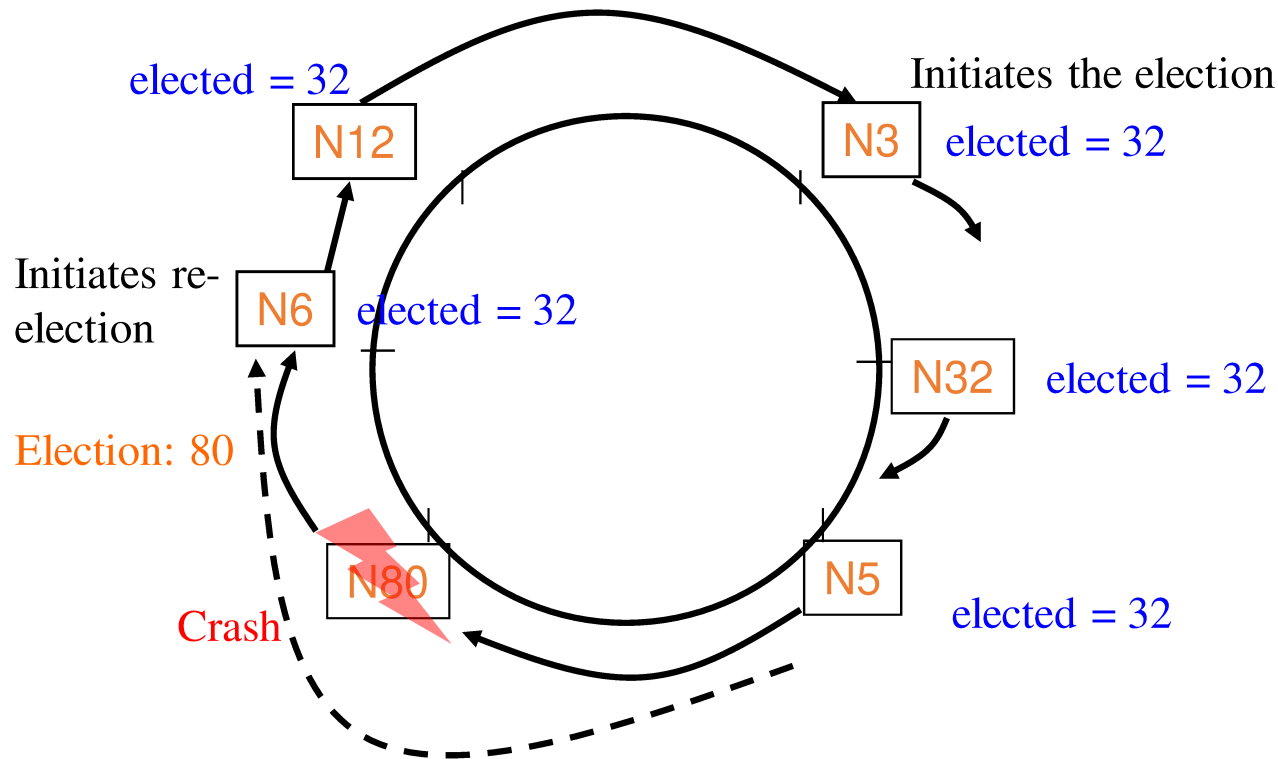
Handling Failures



Handling failures

- Use the failure detector.
- A process can detect failure of N80 via its own local failure detector:
 - Repair the ring.
 - Stop forwarding Election:80 message.
 - Start a new run of leader election.


Handling Failures




Handling failures

- Use the failure detector.
- A process that detects the failure of N80 via its own local failure detector:
 - Repair the ring.
 - Stop forwarding Election:80 message.
 - Start a new run of leader election.
- But failure detectors cannot be both complete and accurate.
 - Incomplete FD => N80's failure might be missed .

What happens if a process failure is undetected?




What happens if a process failure is undetected?




Handling failures

- Use the failure detector.
- A process can detect failure of N80 via its own local failure detector:
 - Repair the ring.
 - Stop forwarding Election:80 message.
 - Start a new run of leader election.
- But failure detectors cannot be both complete and accurate.
 - Incomplete FD => N80's failure might be missed
 - violation of liveness
 - Inaccurate FD => N80 mistakenly detected as failed


What can happen if alive process is detected as failed?




What can happen if alive process is detected as failed?




What can happen if alive process is detected as failed?



What can happen if alive process is detected as failed?



What can happen if an alive process is detected as failed?



Safety has been violated.

Fixing for failures

- Use the failure detector.
- A process can detect failure of N80 via its local failure detector:
 - Repair the ring.
 - Stop forwarding Election:80 message.
 - Start a new run of leader election.
- But failure detectors cannot be both complete and accurate.
 - Incomplete FD => N80's failure might be missed
 - violation of liveness
 - Inaccurate FD => N80 mistakenly detected as failed
 - new ring will be constructed without N80.
 - a process with lower attribute will be selected.
 - violation of safety

Classical Election Algorithms

- Ring election algorithm
- Bully algorithm

Bully algorithm

- Faster **turnaround time** than ring election.
- Explicitly **build in notion of timeouts into the algorithm.**
- Let's assume (for simplicity of exposition) that the attribute based on which leader is elected is the process id.
- Before discussing Bully algorithm, let's first discuss a simpler (related) algorithm.....

Multicast-based algorithm

- Start an election
 - Multicast <election, my ID> to all processes
 - If receive <agree> from all processes, then elected
 - Multicast <coordinator, my ID>
 - If receive <disagree> from any process
 - Give up election
- Receive <election, ID> from process p
 - If ID > my ID
 - Reply <agree> to p (unicast)
 - If ID < my ID
 - Reply <disagree> to p
 - Start election (if not already running)
- What about failures?

Multicast-based algorithm

- Start an election
 - Multicast <election, my ID> to all processes
 - If receive <agree> from all processes or timeout, then elected
 - Multicast <coordinator, my ID>
 - If receive <disagree> from any process
 - Give up election
- Receive <election, ID> from process p
 - If ID > my ID
 - Send <agree> to p (unicast)
 - If ID < my ID
 - Send <disagree> to p
 - Start election (if not already running)
- Can we improve on this?

Bully Algorithm

- All processes know other process' ids.
- Do not need to multicast **election** to all processes.
- Only to processes with higher id.

Bully Algorithm

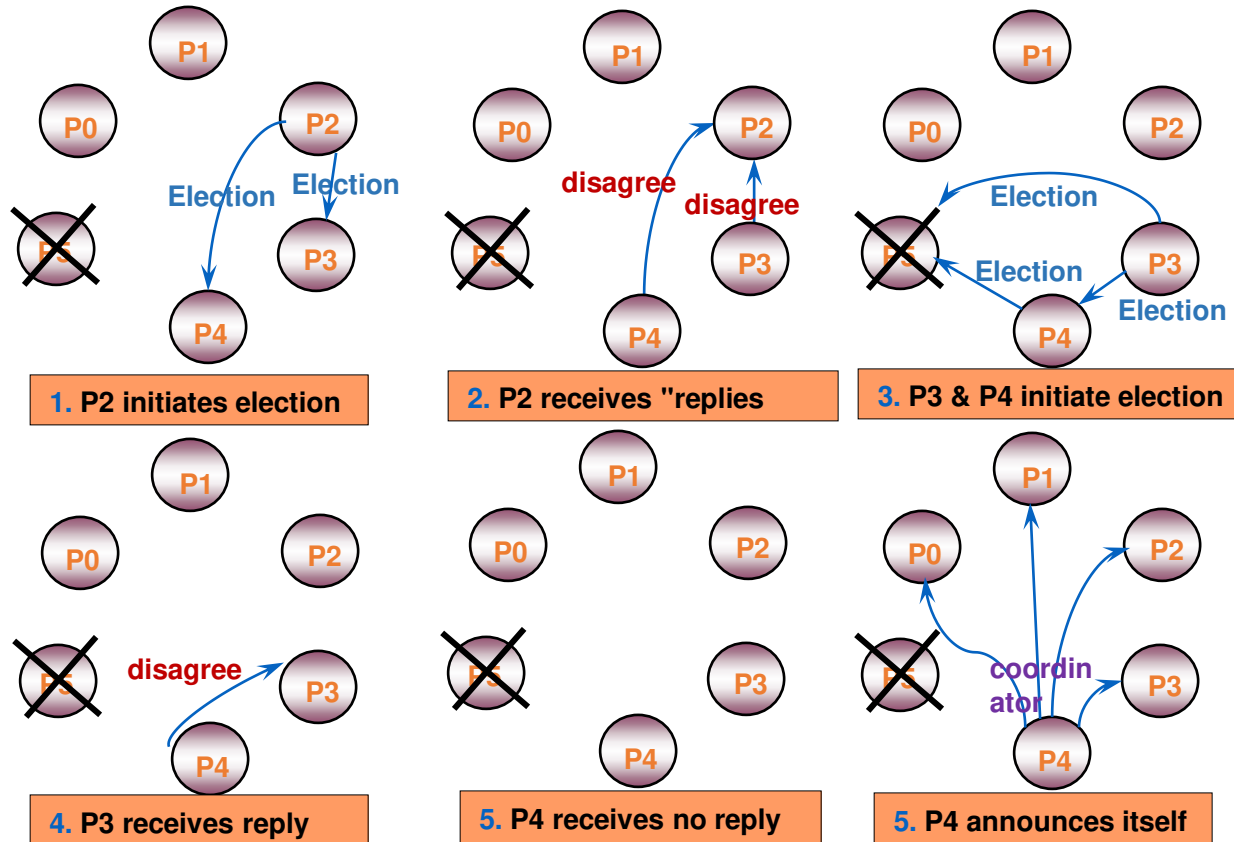
- When a process wants to initiate an election
 - if it knows its id is the highest
 - it elects itself as coordinator, and sends a *Coordinator* message to all processes with lower identifiers.
Election is completed.
 - else
 - it initiates an election by sending an *Election* message
(contd. on next page)

Bully Algorithm (2)

- **else** it initiates an election by sending an *Election* message
 - Sends it to only **processes** that **have a higher id than itself**.
 - **if** receives no answer **within timeout**, **calls itself leader** and sends *Coordinator* message to all **lower id processes**. Election completed.
 - **if** answer received, then there is some non-faulty higher process => so, **wait for coordinator message**. If none received after another timeout, start a new election run.
- A process that receives an *Election* message **replies with disagree message**, and starts its own leader election protocol (unless it has already done so).

Bully Algorithm: Example


P2 initiates election after detecting P5's failure.



What if P4 fails after step 3?

Bully Algorithm: Example

P2 initiates election after detecting P5's failure.



What if P4 fails after step 4?

Bully Algorithm (2)

- **else** it initiates an election by sending an *Election* message
 - Sends it to only processes that have a *higher id than itself*.
 - **if** receives no answer within *timeout*, calls itself leader and sends *Coordinator* message to all lower id processes.
Election completed.
 - **if** answer received however, then there is some non-faulty higher process => so, wait for coordinator message. **If none received after another *timeout*, start a new election run.**
- A process that receives an *Election* message replies with *disagree* message, and starts its own leader election protocol (unless it has already done so).

Timeout values

- Assume the one-way message transmission time (T) is known.
- First timeout value (when the process that has initiated election waits for the first response)
 - Must be set as accurately as possible.
 - If it is too small, a lower id process can declare itself to be the coordinator even when a higher id process is alive.
 - What should be the first timeout value be, given the above assumption?
 - $2T + (\text{processing time}) \approx 2T$
- When the second timeout happens (after 'disagree' message), election is re-started.
 - A very small value will lead to extra "Election" messages.
 - A suitable option is to use the worst-case turnaround time.

Next Class

- Analysis of Bully Algorithm
- Consensus