

Distributed Systems

ECE428

Lecture 15

Adopted from Spring 2021

Agenda for today

- Consensus
 - Consensus in synchronous systems
 - *Chapter 15.4*
 - Impossibility of consensus in asynchronous systems
 - *We will not cover the proof in details*
 - Good enough consensus algorithm for asynchronous systems:
 - *Paxos made simple, Leslie Lamport, 2001*
 - Other forms of consensus algorithm
 - Raft (log-based consensus)
 - Block-chains (distributed consensus)

Raft: A Consensus Algorithm for Replicated Logs

Slides from Diego Ongaro and John Ousterhout, Stanford
University

Raft Overview

1. Leader election:
 - Select one of the servers to act as leader
 - Detect crashes, choose new leader
2. Neutralizing old leaders
3. Normal operation (basic log replication)
4. Safety and consistency after leader changes

Raft Overview

1. Leader election:
 - Select one of the servers to act as leader
 - Detect crashes, choose new leader
2. Neutralizing old leaders
3. Normal operation (basic log replication)
4. Safety and consistency after leader changes

Heartbeats and Timeouts

- Servers start up as **followers**.
- Followers expect to receive RPCs from leaders or candidates.
- Leaders must send **heartbeats** (empty AppendEntries RPCs) to maintain authority.
- If **electionTimeout** elapses with no RPCs:
 - Follower assumes leader has crashed
 - Follower promotes itself to **candidate** and starts new election
 - Timeouts typically in range 100-500ms
 - *Randomly chosen in some range to reduce probability of split election.*

Election Basics

- On timeout:
 - Increment current term
 - Change to Candidate state
 - Vote for self
 - Send RequestVote RPCs to all other servers:
 1. Receive votes from majority of servers:
 - Become leader
 - Send AppendEntries heartbeats (RPCs) periodically to all other servers
 2. Receive RPC from valid leader (with same or higher term):
 - Return to follower state
 3. No-one wins election (election timeout elapses):
 - Increment term, start new election

Election Basics: RequestVote RPCs

- Suppose a server in term `currentTerm` has voted for process with id `votedFor` in that term.
- When it receives RequestVote RPC from process `candidatId` with term `voteRequestTerm`:

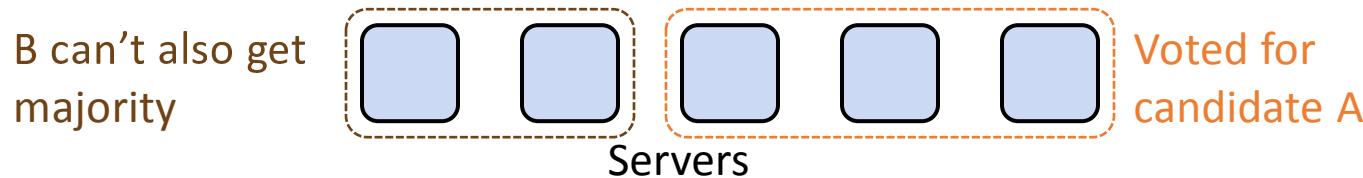
If `voteRequestTerm < currentTerm`
 reply false
 return.

If `voteRequestTerm > currentTerm`
 `currentTerm = voteRequestTerm`, `votedFor = null`

If (`votedFor` is null or `candidatId`)*
 //should not have voted for anyone else in that term
 Grant vote, `votedFor = candidatId`
 **we will extend on this condition later.*

Elections, cont'd

- **Safety:** allow at most one winner per term
 - Each server gives out only one vote per term (persist on disk)
 - Two different candidates can't accumulate majorities *in same term*



- **Liveness:** some candidate must eventually win
 - Choose election timeouts randomly in $[T, kT]$
 - One server usually times out and wins election before others wake up
 - Works well if $T \gg$ broadcast time
- *Safety is guaranteed. Liveness is not guaranteed.*

Implication of terms

- Each term has at most one leader (safety condition).
- Terms always increase with time.
- If the latest term has an elected leader, majority of processes must have updated themselves to the latest term.
- Only the leader can commit log entries (we will discuss this next).

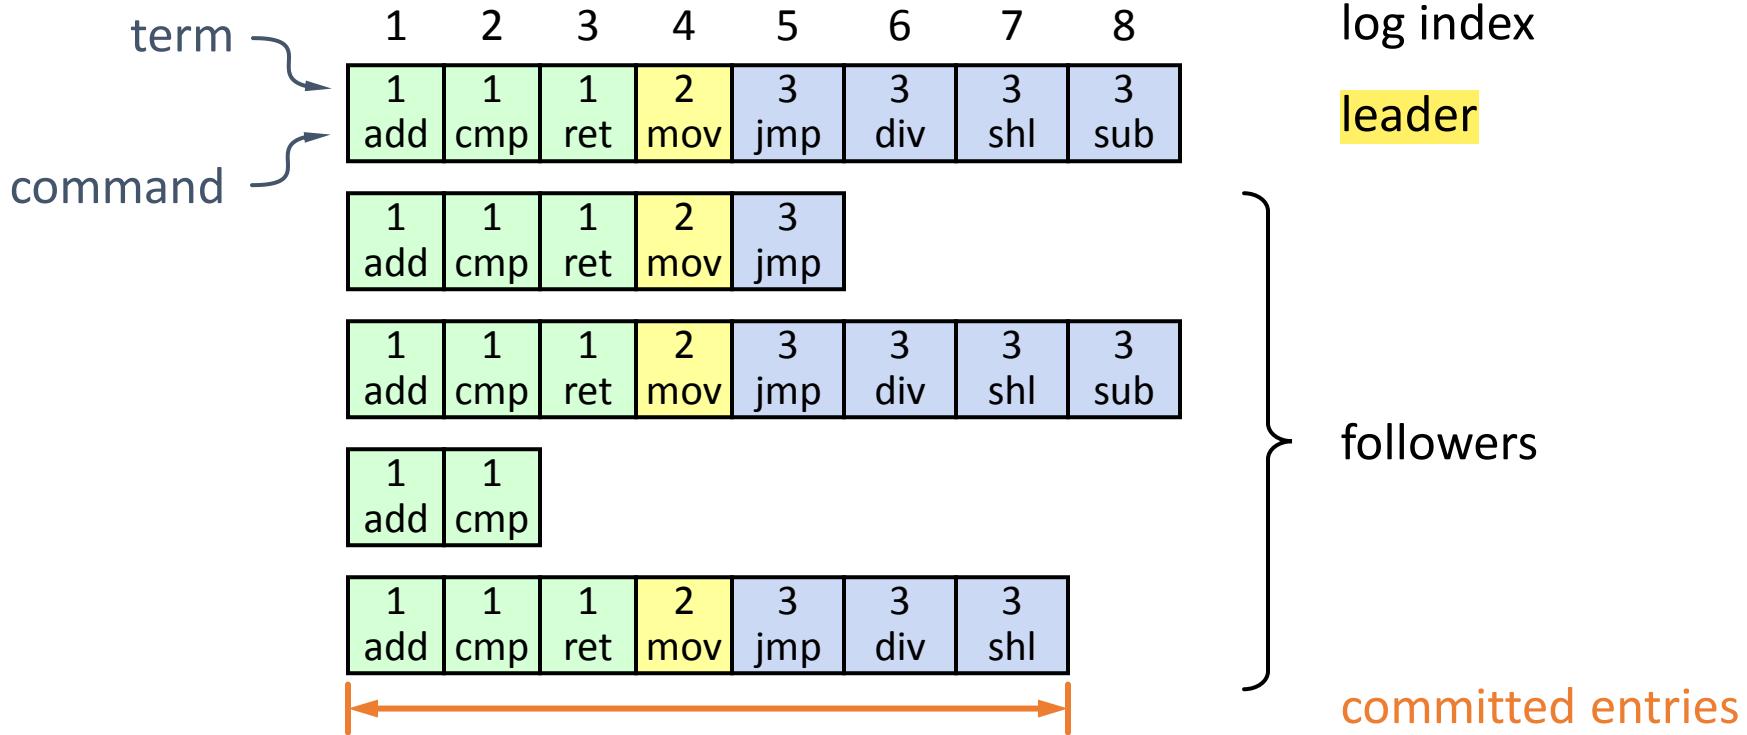
Neutralizing Old Leaders

- Deposed leader may not be dead:
 - Temporarily disconnected from network
 - Other servers elect a new leader
 - Old leader becomes reconnected, attempts to commit log entries
- Terms used to detect stale leaders (and candidates)
 - Every RPC contains term of sender
 - If sender's term is older, RPC is rejected, sender reverts to follower and updates its term
 - If receiver's term is older, it reverts to follower, updates its term, then processes RPC normally
- Election updates terms of majority of servers
 - Deposed server cannot commit new log entries

Raft Overview

1. Leader election:
 - Select one of the servers to act as leader
 - Detect crashes, choose new leader
2. Neutralizing old leaders
3. Normal operation (basic log replication)
4. Safety and consistency after leader changes

Log Structure



- Log entry = **index, term, command**
- Log stored on stable storage (disk); **survives crashes**
- Entry is **committed** by the leader when certain conditions are met.
 - Durable, will eventually be executed by state machines
 - * we will get back to this.

Normal Operation

- Client sends command to leader
- Leader appends command to its log (*not yet committed*)
- Leader sends AppendEntries RPCs to followers
- Once new entry committed* (*we will discuss when and how*):
 - Leader passes command to its state machine, returns result to client
 - Leader notifies followers of committed entries in subsequent AppendEntries RPCs
 - Followers pass committed commands to their state machines
- Crashed/slow followers?
 - Leader retries RPCs until they succeed
- Performance is optimal in common case:
 - One successful RPC to any majority of servers

Log Consistency

High level of coherency between logs:

Raft guarantees that:

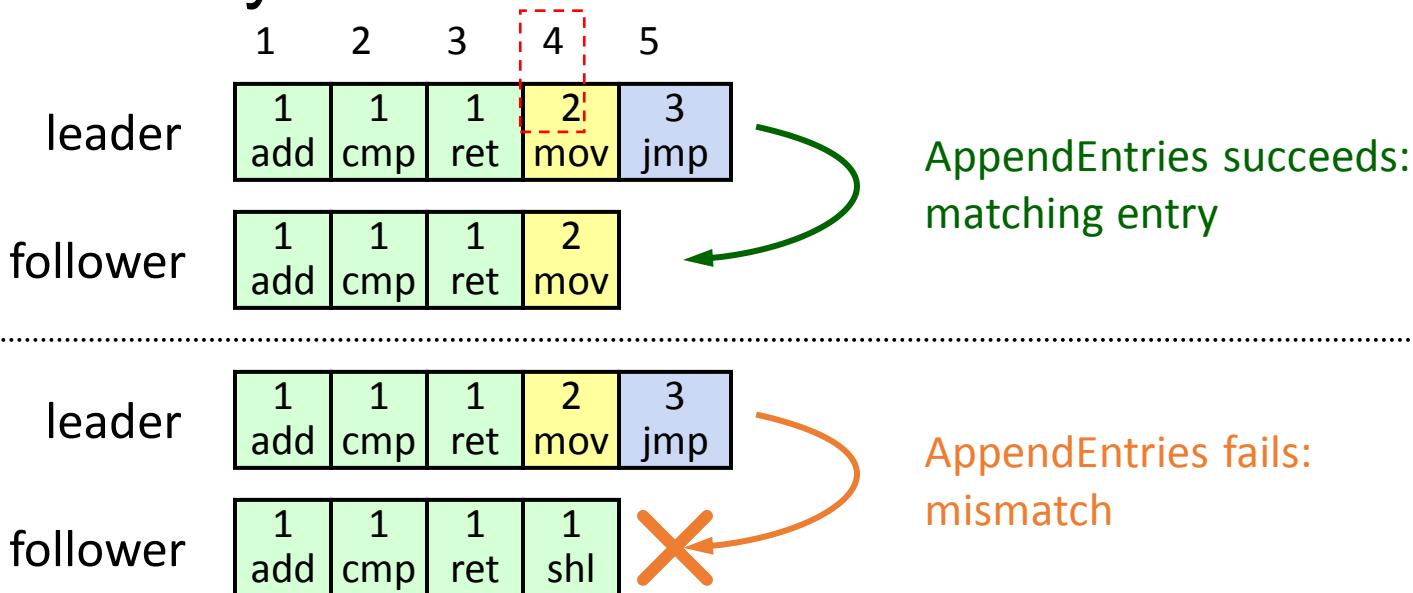
- If log entries on different servers have same index and term:
 - They store the same command
 - The logs are identical in all preceding entries

1	2	3	4	5	6	
1 add	1 cmp	1 ret	2 mov	3 jmp	3 div	
1 add	1 cmp	1 ret	2 mov	3 jmp	4 sub	4 add

- If a given entry is committed, all preceding entries are also committed

AppendEntries Consistency Check

- Each AppendEntries RPC contains **index**, and **term** of **entry preceding new ones**
- Follower must contain matching entry; otherwise it rejects request
- Implements an **induction step**, ensures coherency

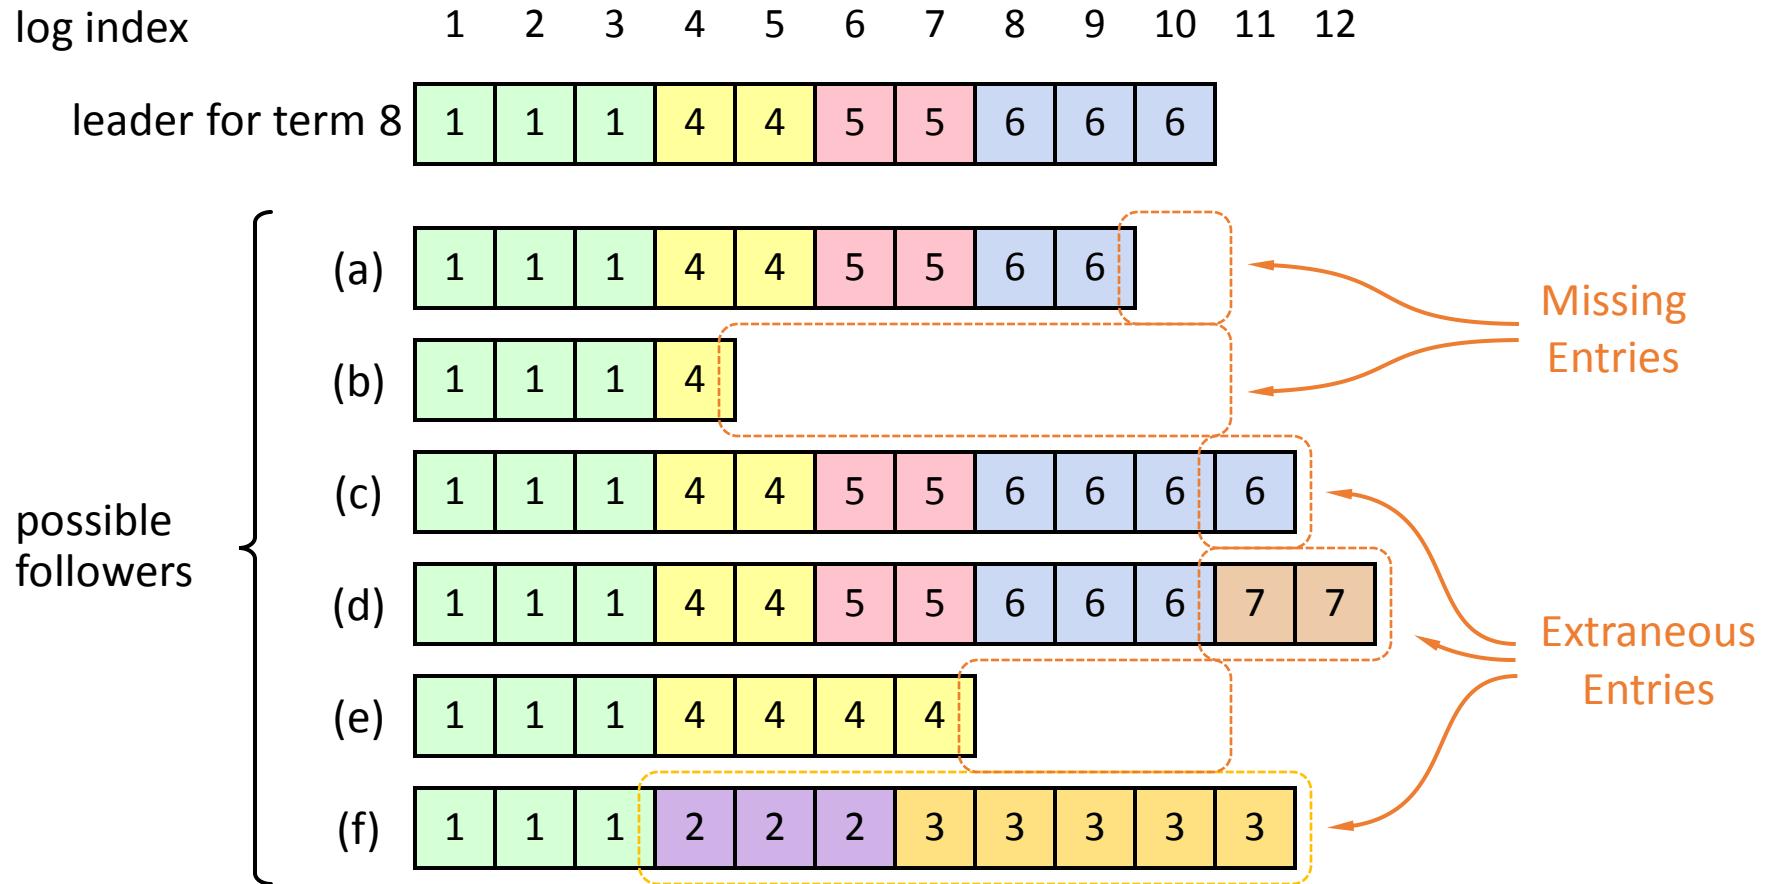


Leader Changes

- At beginning of new leader's term:
 - Old leader may have **left entries partially replicated**
 - No special steps by new leader: just **start normal operation**
 - Leader's log is “the truth”
 - Will eventually make follower's logs identical to leader's
 - *Unless a new leader gets elected during the process.*
 - Multiple crashes can leave many extraneous log entries:

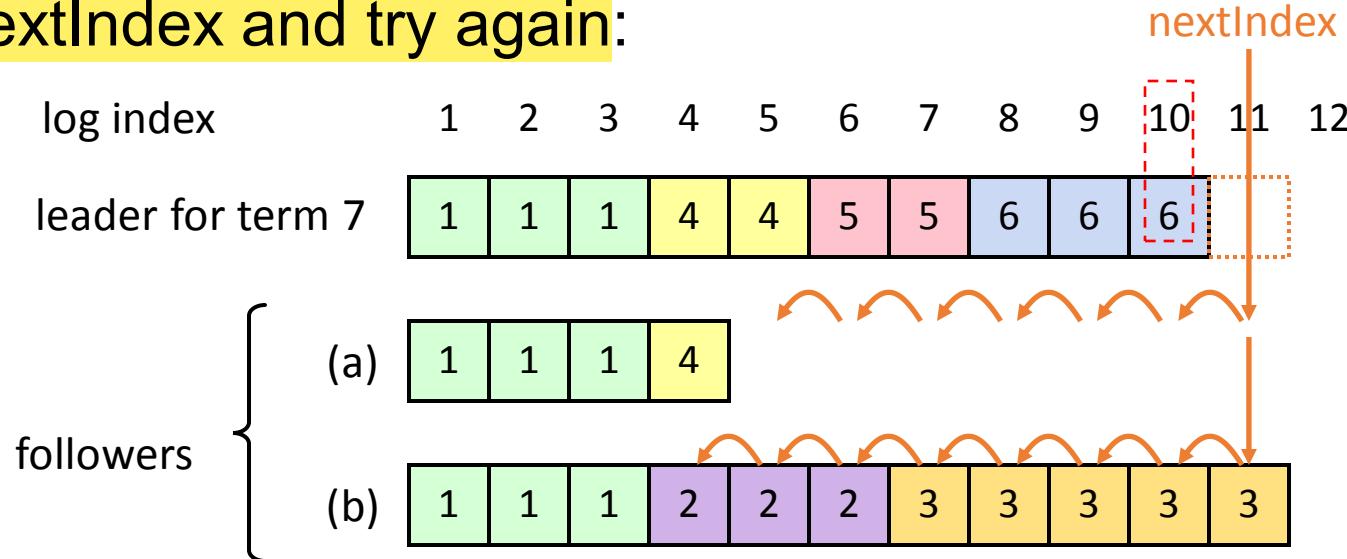
	log index	1	2	3	4	5	6	7	8
term	s ₁	1	1	5	6	6	6		
	s ₂	1	1	5	6	7	7	7	
	s ₃	1	1	5	5				
	s ₄	1	1	2	3				
	s ₅	1	1	2	2	4	4	4	

Log Inconsistencies



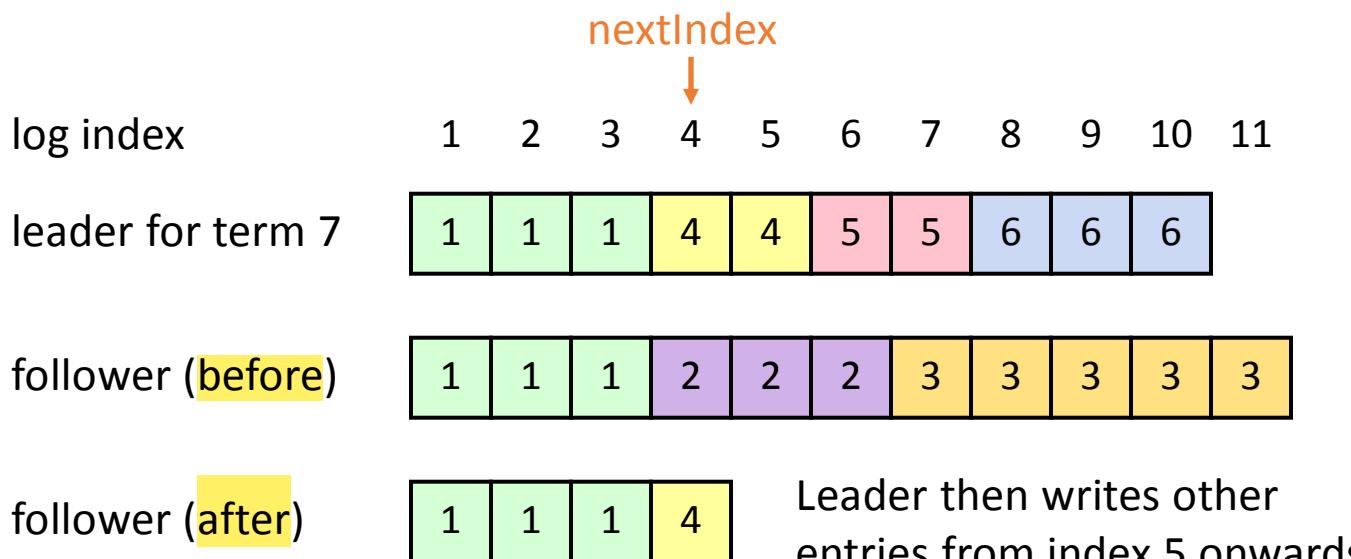
Repairing Follower Logs

- New leader must make follower logs consistent with its own
 - Delete extraneous entries
 - Fill in missing entries
- Leader keeps **nextIndex** for each follower:
 - Index of next log entry to send to that follower
 - Initialized to $(1 + \text{leader's last index})$
- When AppendEntries consistency check fails, decrement **nextIndex** and try again:



Repairing Logs, cont'd

- When follower overwrites inconsistent entry, it deletes all subsequent entries:



Safety Requirement for log consensus

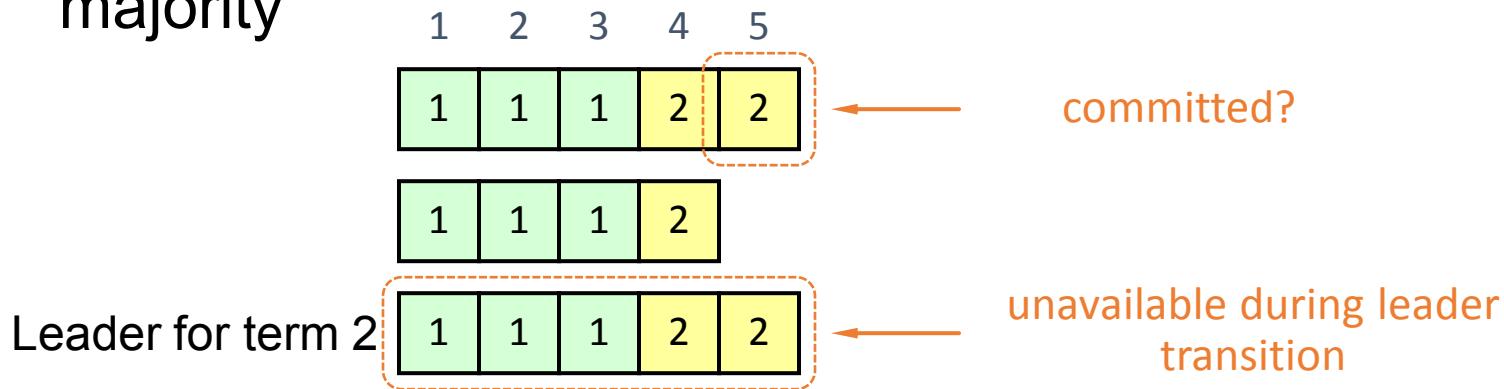
Once a log entry has been applied to a state machine, no other state machine must apply a different value for that log entry.

- Raft safety property:
 - If a leader has decided that a log entry is committed, that entry will be present in the logs of all future leaders
- This guarantees the safety requirement
 - Leaders never overwrite entries in their logs
 - Only entries in the leader's log can be committed
 - Entries must be committed before applying to state machine



Picking the Best Leader

- During elections, choose candidate with log **most likely** to contain all committed entries
 - Candidates include log info in RequestVote RPCs (index & term of last log entry)
 - Voting server V denies vote if its log is “more *up-to-date*”:
 $(\text{lastTerm}_V > \text{lastTerm}_C) \parallel (\text{lastTerm}_V == \text{lastTerm}_C) \&\& (\text{lastIndex}_V > \text{lastIndex}_C)$
 - Leader will have “most complete” log among electing majority



Election Basics: RequestVote RPCs

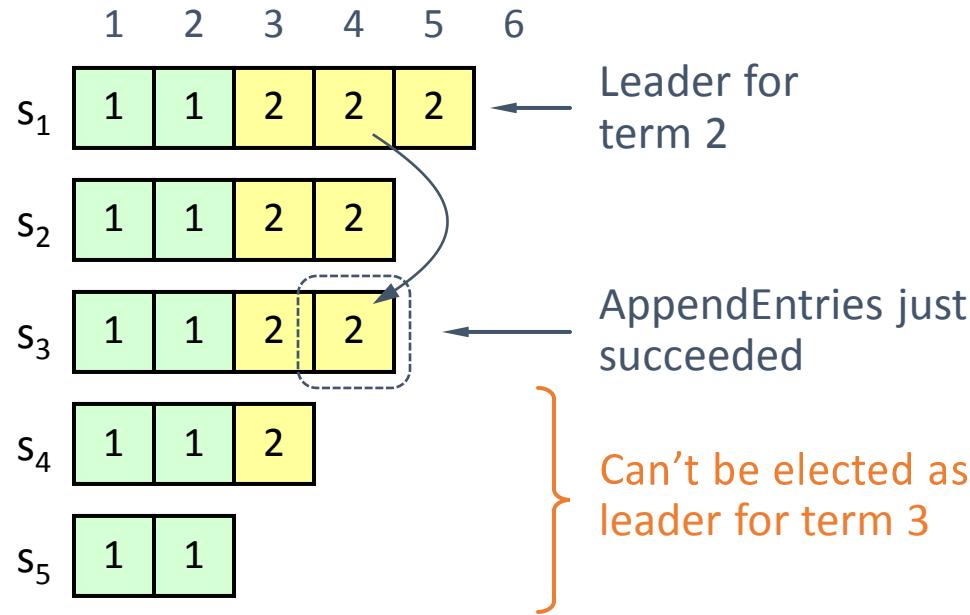
- Suppose a server **S** in term **currentTerm** has voted for process with id **votedFor** in that term.
- When it receives RequestVote RPC from process **candidateId** with term **voteRequestTerm**:
 - If **voteRequestTerm < currentTerm**
 Reply false, return.
 - If **voteRequestTerm > currentTerm**
 currentTerm = voteRequestTerm, **votedFor = null**
 If (candidate's log is at least as *up-to-date* S's log) and (**votedFor** is null or **candidateId**)
 Grant vote, **votedFor = candidateId**

Checking for up-to-date logs

- Log L1 is **more up-to-date** than Log L2 iff:
 - term of the last entry in L1 > term of the last entry in L2, or
 - (term of the last entry in L1 = term of the last entry in L2) and (length of L1 is larger than length of L2).
- Can check length of log by comparing the last entry's index.
- So RequestVote RPC argument also contains the candidate's last log entry's index and term

Committing Entry from Current Term

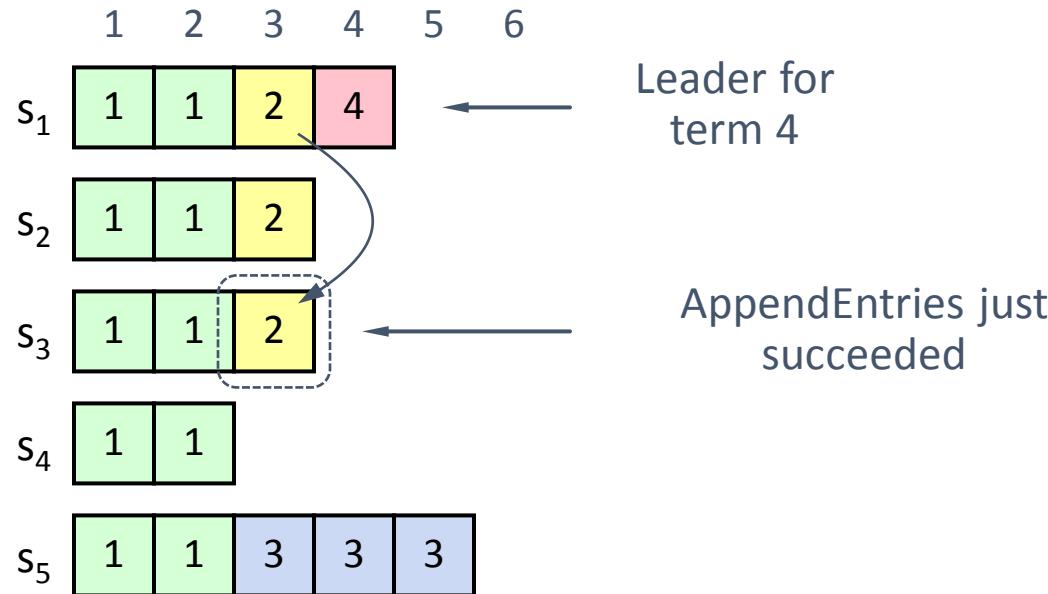
- *When can a leader commit entries?*



- Leader decides entry in term 4 is committed
 - Safe: leader for term 3 must contain entry 4
- What about committing entry in term 5?
- *Perhaps leader can commit an entry once replicated on majority of servers?*

Committing Entry from Earlier Term

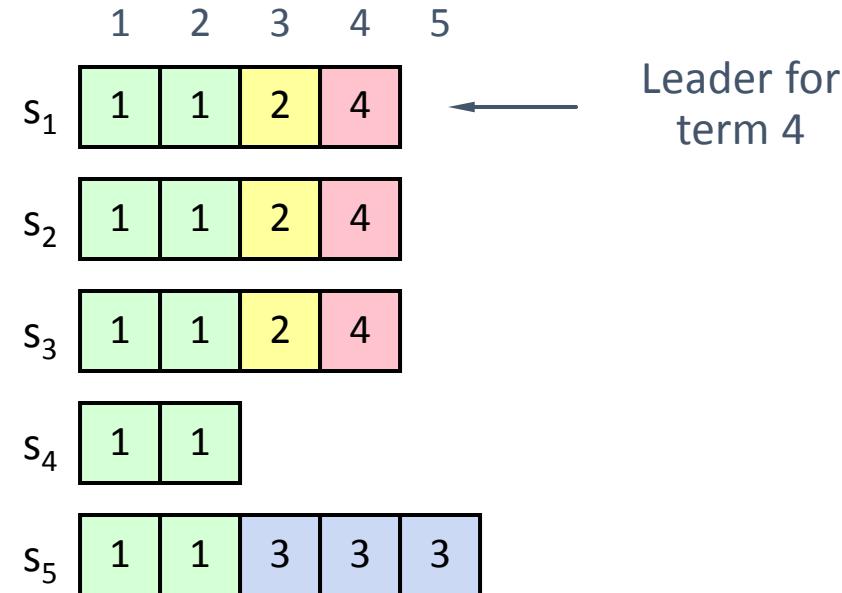
- Leader is trying to finish committing entry from an earlier term



- Entry 3 **not safely committed**:
 - s₅ can be elected as leader for term 5
 - If elected, it will overwrite entry 3 on s₁, s₂, and s₃!

New Commitment Rules

- For a leader to decide an entry is committed:
 - Must be stored on a majority of servers
 - *At least one new entry from leader's current term must also be stored on majority of servers*
- Once entry 4 committed:
 - s_5 cannot be elected leader for term 5
 - Entries 3 and 4 both safe



Combination of election rules and commitment rules makes Raft safe.

Raft Protocol Summary

Followers

- Respond to RPCs from candidates and leaders.
- Convert to candidate if election timeout elapses without either:
 - Receiving valid AppendEntries RPC, or
 - Granting vote to candidate

Candidates

- Increment currentTerm, vote for self
- Reset election timeout
- Send RequestVote RPCs to all other servers, wait for either:
 - Votes received from majority of servers: become leader
 - AppendEntries RPC received from new leader: step down
 - Election timeout elapses without election resolution: increment term, start new election
 - Discover higher term: step down

Leaders

- Initialize nextIndex for each to last log index + 1
- Send initial empty AppendEntries RPCs (heartbeat) to each follower; repeat during idle periods to prevent election timeouts
- Accept commands from clients, append new entries to local log
- Whenever last log index \geq nextIndex for a follower, send AppendEntries RPC with log entries starting at nextIndex, update nextIndex if successful
- If AppendEntries fails because of log inconsistency, decrement nextIndex and retry
- Mark log entries committed if stored on a majority of servers and at least one entry from current term is stored on a majority of servers
- Step down if currentTerm changes

Persistent State

Each server persists the following to stable storage synchronously before responding to RPCs:

currentTerm	latest term server has seen (initialized to 0 on first boot)
votedFor	candidateId that received vote in current term (or null if none)
log[]	log entries

Log Entry

term	term when entry was received by leader
index	position of entry in the log
command	command for state machine

RequestVote RPC

Invoked by candidates to gather votes.

Arguments:

candidateId	candidate requesting vote
term	candidate's term
lastLogIndex	index of candidate's last log entry
lastLogTerm	term of candidate's last log entry

Results:

term	currentTerm, for candidate to update itself
voteGranted	true means candidate received vote

Implementation:

1. If term > currentTerm, currentTerm \leftarrow term
(step down if leader or candidate)
2. If term == currentTerm, votedFor is null or candidateId, and candidate's log is at least as complete as local log, grant vote and reset election timeout

AppendEntries RPC

Invoked by leader to replicate log entries and discover inconsistencies; also used as heartbeat .

Arguments:

term	leader's term
leaderId	so follower can redirect clients
prevLogIndex	index of log entry immediately preceding new ones
prevLogTerm	term of prevLogIndex entry
entries[]	log entries to store (empty for heartbeat)
commitIndex	last entry known to be committed

Results:

term	currentTerm, for leader to update itself
success	true if follower contained entry matching prevLogIndex and prevLogTerm

Implementation:

1. Return if term < currentTerm
2. If term > currentTerm, currentTerm \leftarrow term
3. If candidate or leader, step down
4. Reset election timeout
5. Return failure if log doesn't contain an entry at prevLogIndex whose term matches prevLogTerm
6. If existing entries conflict with new entries, delete all existing entries starting with first conflicting entry
7. Append any new entries not already in the log
8. Advance state machine with newly committed entries

More details in Raft paper

- Link on the course website:
<https://raft.github.io/raft.pdf>
- Check also Raft tutorials on Blackboard
- The concepts covered Section 6 and beyond
are not in your syllabus.