

Distributed Systems

ECE428

Lecture 17

Adopted from Spring 2021

Agenda for the next 2-3 classes

- Transaction Processing and Concurrency Control
 - Chapter 16
 - Transaction semantics: ACID
 - Isolation and serial equivalence
 - Conflicting operations
 - Two-phase locking
 - Deadlocks
 - Timestamped ordering
- First focus on transactions executed on a single server.
- Look into distributed transactions later (Chapter 17)

Today's Agenda

- Transaction Processing and Concurrency Control
 - Chapter 16
 - Transaction semantics: ACID
 - Isolation and serial equivalence
 - Conflicting operations
 - Two-phase locking
 - Deadlocks
 - Timestamped ordering
- First focus on transactions executed on a single server.
- Look into distributed transactions later (Chapter 17)

Transaction

- Series of *operations* executed by a client on a server (or a set of servers).
- Example: Switch from T4 to T3 section:
rosters.remove("ece428", "t4", student.name)
student.schedule.remove("ece428", "t4")
student.schedule.add("ece428", "t3")
rosters.add("ece428", "t3", student.name)

Transaction

- Another example:

Client code:

```
int transaction_id = openTransaction();  
x = server.getFlightAvailability(ABC123, date); // read(ABC123, date)  
if (x > 0)  
    y = server.bookTicket(ABC123, date); // write(ABC123, date)  
server.putSeat(y, "aisle");  
closeTransaction(transaction_id);
```

Transaction Properties

- Atomic: all-or-nothing
 - Transaction either executes completely or not at all
- Consistent: required rules are maintained
- Isolation: multiple transactions do not interfere with each other
 - Equivalent to running transactions in isolation
- Durability: values preserved even after crashes

Transaction Properties

- Atomic: all-or-nothing
 - Transaction either executes completely or not at all
- Consistent: rules maintained
- Isolation: multiple transactions do not interfere with each other
 - Equivalent to running transactions in isolation
- Durability: values preserved even after crashes

ACID properties

Atomicity

- All-or-nothing
 - Transaction either executes completely or not at all
- What can happen after partial execution?

```
rosters.remove("ece428", "t4", student.name)  
student.schedule.remove("ece428", "t4")  
student.schedule.add("ece428", "t3")  
rosters.add("ece428", "t3", student.name)
```


Atomicity

- All-or-nothing
 - Transaction either executes completely or not at all
- Make tentative updates to data.
- **Commit** transaction to make tentative updates permanent.
- **Abort transaction** to **roll back** to previous values.

Consistency

Various rules about state of objects must be maintained:

- E.g. class enrollment limit, schedule can't conflict
- Account balances have to stay positive
- Consistency must be maintained at end of transaction.
- Checked at commit time, abort if not satisfied

```
rosters.remove("ece428", "t4", student.name)  
student.schedule.remove("ece428", "t4")  
student.schedule.add("ece428", "t3")  
rosters.add("ece428", "t3", student.name)
```

Durability

- Committed transactions must persist:
 - Client **crashes**
 - Server crashes
- How do we ensure this?
 - **Replication**
 - Permanent storage

Isolation

Multiple clients may execute transactions concurrently on one server.

What could go wrong?

What could go wrong?

Transaction T1

```
x = getSeats(ABC123);
```

```
// x = 10
```

```
if(x > 0)
```

```
    x = x - 1;
```

```
write(x, ABC123);
```

```
commit
```

Transaction T2

```
x = getSeats(ABC123);
```

```
if(x > 0) // x = 10
```

```
    x = x - 1;
```

```
write(x, ABC123);
```

```
commit
```

At Server: seats = 10

seats = 9

seats = 9

1. Lost Update Problem

Transaction T1

x = getSeats(ABC123);

// x = 10

if(x > 0)

x = x - 1;

write(x, ABC123);

commit

Transaction T2

x = getSeats(ABC123);

if(x > 0) // x = 10

x = x - 1;

write(x, ABC123);

commit

At Server: seats = 10

seats = 9

seats = 9

T1's or T2's update was lost!

What else could go wrong?

Transaction T1

```
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(x-5, ABC123);
```

// ABC123 = 5 now

```
write(y+5, ABC789);
```

```
commit
```

Transaction T2

```
x = getSeats(ABC123);  
y = getSeats(ABC789);
```

// x = 5, y = 15

```
print("Total:" x+y);
```

// Prints "Total: 20"

```
commit
```

At Server:

ABC123 = 10

ABC789 = 15

2. Inconsistent Retrieval Problem

Transaction T1

```
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(x-5, ABC123);
```

// ABC123 = 5 now

```
write(y+5, ABC789);
```

```
commit
```

Transaction T2

```
x = getSeats(ABC123);  
y = getSeats(ABC789);
```

// x = 5, y = 15

```
print("Total:" x+y);  
// Prints "Total: 20"
```

```
commit
```

At Server:

ABC123 = 10

ABC789 = 15

**T2's sum is the
wrong value!
Should have been
"Total: 25"**

Isolation

Multiple clients executing transactions concurrently on one server.

What could go wrong?

- Lost Update Problem
- Inconsistent Retrieval Problem

How to prevent transactions from affecting each other?

Isolation

How to prevent transactions from affecting each other?

- Option 1: Execute them serially at the server (one at a time).
 - Grab a global lock before executing any transaction, release the lock after the transaction has committed (or aborted).
- But this reduces the number of concurrent transactions
 - Transactions per second often directly related to revenue of the companies
 - This metric needs to be maximized

Goal: increase concurrency while maintaining correctness (ACID).

Concurrent Transactions

Goal: increase concurrency while maintaining correctness (ACID).

- How do we increase concurrency?
 - Instead of targeting strict serial execution, target serial equivalence.

Interleaving

- An ordered sequence of the operations across multiple transactions, where each transaction's operations follows the order defined by the transaction.
- E.g., if $T_1 = \{op_1, op_2, op_3\}$ and $T_2 = \{op^4, op^5, op^6\}$ then $O = \{op_1, op_2, op^4, op_3, op^5, op^6\}$ is an example of interleaving.

Interleaving: Example

Transaction T1

```
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(x-5, ABC123);
```

```
write(y+5, ABC789);
```

```
commit
```

Transaction T2

```
x = getSeats(ABC123);  
y = getSeats(ABC789);
```

```
print("Total:" x+y);
```

```
commit
```

Interleaving

```
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(x-5, ABC123);  
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(y+5, ABC789);
```

```
print("Total:" x+y);  
commit  
commit
```

Concurrent Transactions

- Allowing transaction operations to be interleaved with one-another increases concurrency.
- To avoid transactions from affecting one another, the interleaving of operations across transactions must be *serially equivalent*.

Serial Equivalence

- An interleaving (say O) of transaction operations is **serially equivalent** iff (if and only if):
 - There is an ordering (O') of those transactions, one at a time,
 - where the operations of each transaction occur consecutively (in a batch)
 - which gives the **same end-result** (for all objects and transactions) as the original interleaving O
- It says: Cannot distinguish end-result of transaction order O from another serial transaction order O'
- E.g., if $T_1 = \{op_1, op_2, op_3\}$ and $T_2 = \{op^4, op^5, op^6\}$
then $O = \{op_1, op_2, op^4, op_3, op^5, op^6\}$ **is serially equivalent**, if:
 - **end result of O is same as** $\{op_1, op_2, op_3, op^4, op^5, op^6\}$
 - Or end result of O is same as $\{op^4, op^5, op^6, op_1, op_2, op_3, \}$

1. Lost Update Problem

Transaction T1

x = getSeats(ABC123);

// x = 10

if(x > 0)

x = x - 1;

write(x, ABC123);

commit

Transaction T2

x = getSeats(ABC123);

if(x > 0) // x = 10

x = x - 1;

write(x, ABC123);

commit

At Server: seats = 10

seats = 9

seats = 9

T1's or T2's update was lost!

Not serially equivalent.

2. Inconsistent Retrieval Problem

Transaction T1

```
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(x-5, ABC123);
```

// ABC123 = 5 now

```
write(y+5, ABC789);
```

```
commit
```

Transaction T2

```
x = getSeats(ABC123);  
y = getSeats(ABC789);
```

// x = 5, y = 15

```
print("Total:" x+y);
```

// Prints "Total: 20"

```
commit
```

At Server:

ABC123 = 10

ABC789 = 15

**T2's sum is the
wrong value!
Should have been
"Total: 25"**

Not serially equivalent.

Checking for Serial Equivalence

- An operation has an **effect** on
 - The **server** object if it is a **write**
 - The **client** (returned value) if it is a **read**
- Two operations are said to be conflicting operations, if their **combined effect** depends on the **order** they are executed
 - read(x) and write(x): conflicting
 - write(x) and read(x): conflicting
 - write(x) and write(x): conflicting
 - read(x) and read(x): **NOT conflicting**
 - swapping them doesn't change their effects
 - **read/write(x)** and read/write(y): NOT conflicting
 - ok to swap them as **they access different objects.**

Checking for Serial Equivalence (cont.)

- *Two transactions are serially equivalent if and only if all pairs of conflicting operations (the pairs containing one operation from each transaction) are executed in the same order (transaction order) for all objects (data) they both access.*
 - Take all pairs of conflicting operations, one from T1 and one from T2
 - If the T1 operation was reflected first on the server, mark the pair as “(T1, T2)”, otherwise mark it as “(T2, T1)”
 - All pairs should be marked as either “(T1, T2)” or all pairs should be marked as “(T2, T1)”.

1. Lost Update Problem

Transaction T1

`x = getSeats(ABC123);`

`// x = 10`

`if(x > 1)`

`x = x - 1;`

`write(x, ABC123);`

`commit`

Transaction T2

`x = getSeats(ABC123);`

`if(x > 1)`

`// x = 10`

(T1, T2)

`x = x - 1;`

`write(x, ABC123);`

`commit`

At Server: seats = 10

seats = 9

seats = 9

(T2, T1)

(T1, T2)

**Same transaction order not maintained across conflicting operations.
Not serially equivalent.**

2. Inconsistent Retrieval Problem

Transaction T1

```
x = getSeats(ABC123);  
y = getSeats(ABC789);  
write(x-5, ABC123);
```

// ABC123 = 5 now

```
write(y+5, ABC789);
```

```
commit
```

Transaction T2

```
x = getSeats(ABC123);  
y = getSeats(ABC789);
```

// x = 5, y = 15

```
print("Total:" x+y);
```

// Prints "Total: 20"

```
commit
```

At Server:

ABC123 = 10

ABC789 = 15

**Same transaction order not maintained across conflicting operations.
Not serially equivalent.**

How do we handle such conflicts?

- Option 2:
 - At **commit point** of a transaction T, check for **serial equivalence** with all other transactions
 - The transactions that **overlapped in time with T**
 - If not serially equivalent
 - **Abort T**
 - Roll back **(undo) any writes that T did** to server objects
- Aborting all such transactions => wasted work.
 - Can we do better?
 - Can we prevent violations from occurring?

Two Approaches

- Preventing isolation from being violated can be done in two ways
 1. **Pessimistic** concurrency control
 2. **Optimistic** concurrency control

Pessimistic vs. Optimistic

- **Pessimistic**: assume the worst, prevent transactions from accessing the same object
 - E.g., Locking
- **Optimistic**: assume the best, allow transactions to write, but check later
 - E.g., Check at commit time

Pessimistic vs. Optimistic

- **Pessimistic**: assume the worst, prevent transactions from accessing the same object
 - E.g., Locking
- Optimistic: assume the best, allow transactions to write, but check later
 - E.g., Check at commit time

Pessimistic: Exclusive Locking

- Grabbing a global lock is wasteful
 - what if no two transactions access the same object?
- Each object has a lock
 - At most one transaction can be inside lock
 - Before reading or writing object O, transaction T must call `lock(O)`
 - Blocks if another transaction already inside lock
 - After entering lock T can read and write O multiple times
 - When done (or at commit point), T calls `unlock(O)`
 - If other transactions waiting at `lock(O)`, allow one of them in
- Sounds familiar?
 - This is **Mutual Exclusion!**

Can we improve concurrency?

- More concurrency => more transactions per second => more revenue (\$\$\$)
- Real-life workloads have a lot of read-only or read-mostly transactions
 - Exclusive per-object locking reduces concurrency
 - Ok to allow two transactions to concurrently read an object, since read-read is not a conflicting pair

Another approach: Read-Write Locks

- Each object has a lock that can be held in one of two modes
 - **Read mode**: multiple transactions allowed in
 - **Write mode**: exclusive lock
- Before first reading O, transaction T calls **read_lock(O)**
 - T allowed in (does not wait on the lock) only if *all* transactions inside lock for O **all entered via read mode**
 - Not allowed (i.e. must wait) if *any* transaction inside lock for O entered **via write** mode

Read Locks Example

read_lock(A)
read(A)

Allowed!
read_lock(A)
read(A)

write_lock(A)
write(A)

Blocked!
read_lock(A)
read(A)

Another approach: Read-Write Locks

- Before first writing O, call `write_lock(O)`
 - Allowed in only if no other transaction inside lock
- If T already holds `read_lock(O)`, and wants to write, call `write_lock(O)` to *promote* lock from read to write mode
 - Succeeds only if no other transactions in write mode or read mode
 - Otherwise, T blocks
- `Unlock(O)` called by transaction T releases any lock on O by T

Write Locks Example

read_lock(A)
read(A)

Blocked!
write_lock(A)
write(A)

write_lock(A)
write(A)

Blocked!
write_lock(A)
write(A)

Write Locks Example

Within a single transaction

read_lock(A)

read(A)

Promoted and allowed!

write_lock(A)

write(A)

Write Locks Example

read_lock(A)
read(A)

Allowed!
read_lock(A)
read(A)

Blocked!
write_lock(A)
write(A)

When to release locks?

- We can have **per-object locks in two modes** to increase concurrency.
- Grab the object's lock in the appropriate mode when trying to access an object.
- *When to release locks?*

write_lock(A)
write(A)
unlock(A)

write_lock(A)
write(A)
unlock(A)
read_lock(A)
read(A)
unlock(A)

Is this a good idea?

Next Class

- Pessimistic Concurrency Control
 - Two-phase Locking
 - Deadlocks
- Optimistic Concurrency Control